

# DA6401:Introduction to Deep Learning

## Module 1B-Modern Neuron & MLP

Ganapathy Krishnamurthi

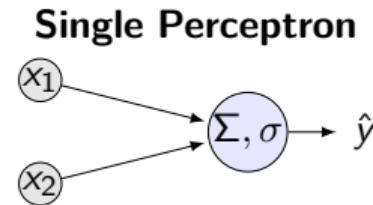
Department of Data Science and Artificial Intelligence , IIT Madras

# The Challenge: Beyond the Single Perceptron

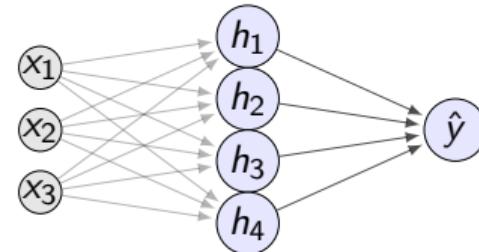
- **Context:** Single perceptrons cannot solve non-linear problems (e.g., XOR).

# The Challenge: Beyond the Single Perceptron

- **Context:** Single perceptrons cannot solve non-linear problems (e.g., XOR).
- **Solution:** Networks of neurons (Multi-Layer Perceptrons) can represent non-linear decision boundaries.



**Multi-Layer  
Perceptron**



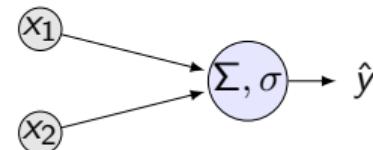
# The Challenge: Beyond the Single Perceptron

- **Context:** Single perceptrons cannot solve non-linear problems (e.g., XOR).
- **Solution:** Networks of neurons (Multi-Layer Perceptrons) can represent non-linear decision boundaries.

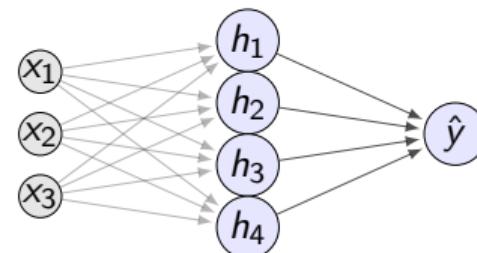
## Learning Problem

We have a learning rule for a single perceptron. But how do we adjust weights deep inside the network?

**Single Perceptron**



**Multi-Layer  
Perceptron**

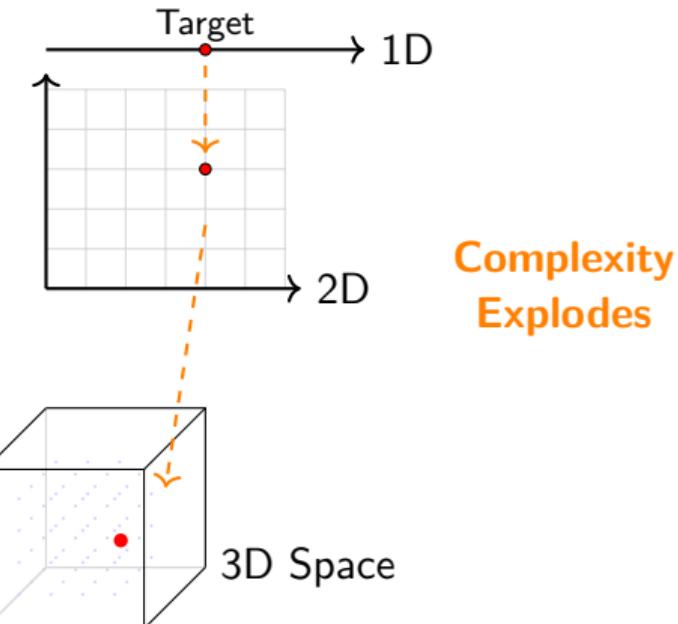


# The Parameter Space Explosion

## The Concept:

- In 1 dimension (1 weight), you just search a line. Easy.
- In 2 dimensions, you search a square. Harder.
- In 3 dimensions, you search a cube.
- **Neural Networks** often have  $10^3$  to  $10^9$  dimensions (weights).

The volume of the space where the "correct" configuration lives becomes infinitesimally small relative to the total volume.



# Approach 1: Stochastic Weight Guessing

- **Concept:** Treat the network as a black box. Randomly sample weight vectors  $\mathbf{w}$  from a high-dimensional space until performance is satisfactory.

# Approach 1: Stochastic Weight Guessing

- **Concept:** Treat the network as a black box. Randomly sample weight vectors  $\mathbf{w}$  from a high-dimensional space until performance is satisfactory.
- **Formal Failure Analysis (Curse of Dimensionality):**
  - Consider a small network with  $N = 1000$  weights.
  - Assume each weight is quantized to just two values,  $w_i \in \{-1, 1\}$ .
  - The search space size is  $|\mathcal{W}| = 2^{1000} \approx 10^{301}$ .
  - For context, the number of atoms in the observable universe is  $\approx 10^{80}$ .

# Approach 1: Stochastic Weight Guessing

- **Concept:** Treat the network as a black box. Randomly sample weight vectors  $\mathbf{w}$  from a high-dimensional space until performance is satisfactory.
- **Formal Failure Analysis (Curse of Dimensionality):**
  - Consider a small network with  $N = 1000$  weights.
  - Assume each weight is quantized to just two values,  $w_i \in \{-1, 1\}$ .
  - The search space size is  $|\mathcal{W}| = 2^{1000} \approx 10^{301}$ .
  - For context, the number of atoms in the observable universe is  $\approx 10^{80}$ .
- **Conclusion:** The volume of the parameter space grows exponentially with the number of parameters. brute-force or undirected random search is theoretically non-viable for practical networks.
- We need a systematic method to navigate this space.

## Approach 2: Parameterization

**Core Idea:** We define our network as a function with a set of learnable parameters (weights and biases), denoted by  $w$ . The goal is to find the optimal values for  $\theta$ .

## Approach 2: Parameterization

**Core Idea:** We define our network as a function with a set of learnable parameters (weights and biases), denoted by  $w$ . The goal is to find the optimal values for  $\theta$ .

### The Network as a Function

We express the model's prediction as

$$\hat{y} = f(x; w),$$

where:

- $x$  is the numerical input data.
- $f$  is the network architecture (e.g., layers, activation functions).
- $w$  represents all the weights and biases in the network.
- $\hat{y}$  is the model's prediction.

# Quantifying Performance: The Loss Function

- To improve, we must quantify "badness." We define a **Loss Function** (or Cost/Objective Function),  $\mathcal{L}$ .

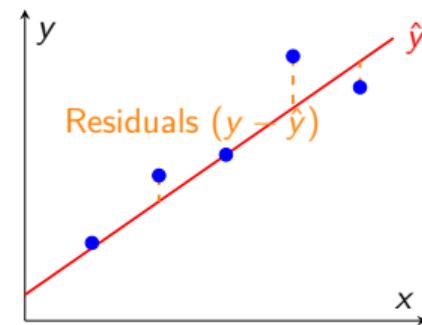
# Quantifying Performance: The Loss Function

- To improve, we must quantify "badness." We define a **Loss Function** (or Cost/Objective Function),  $\mathcal{L}$ .
- Given a dataset  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ , we measure the discrepancy between predictions  $\hat{y}^{(i)} = f(x^{(i)}; \mathbf{w})$  and targets  $y^{(i)}$ .

# Quantifying Performance: The Loss Function

- To improve, we must quantify "badness." We define a **Loss Function** (or Cost/Objective Function),  $\mathcal{L}$ .
- Given a dataset  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ , we measure the discrepancy between predictions  $\hat{y}^{(i)} = f(x^{(i)}; \mathbf{w})$  and targets  $y^{(i)}$ .
- Example: Mean Squared Error (MSE)** used for regression:

$$J(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y^{(i)} - \hat{y}^{(i)})^2$$



**Figure:** Visualizing MSE: The loss function minimizes the sum of squared vertical residuals (orange lines).

# The Parameter Space & Loss Landscape

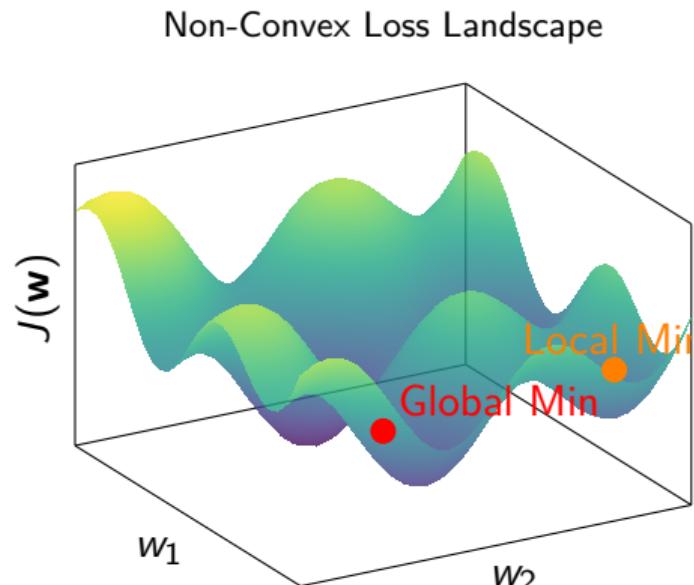
- **Shift in Perspective:** During training, the data  $\mathcal{D}$  is fixed. The network's parameters  $\mathbf{w}$  are the variables.

# The Parameter Space & Loss Landscape

- **Shift in Perspective:** During training, the data  $\mathcal{D}$  is fixed. The network's parameters  $\mathbf{w}$  are the variables.
- We view the loss as a function solely of the weights:  $J(\mathbf{w}) : \mathbb{R}^d \rightarrow \mathbb{R}$ .

# The Parameter Space & Loss Landscape

- **Shift in Perspective:** During training, the data  $\mathcal{D}$  is fixed. The network's parameters  $\mathbf{w}$  are the variables.
- We view the loss as a function solely of the weights:  $J(\mathbf{w}) : \mathbb{R}^d \rightarrow \mathbb{R}$ .
- This defines a high-dimensional "loss landscape."
  - **Coordinates:** The values of the weights  $(w_1, w_2, \dots, w_d)$ .
  - **Altitude:** The loss value  $J(\mathbf{w})$ .



**Figure:** A visualization of a complex 2D weight space resulting in a rugged loss surface.

# Optimization via Calculus

- Assuming  $J(\mathbf{w})$  is differentiable, calculus provides the direction of steepest ascent.

# Optimization via Calculus

- Assuming  $J(\mathbf{w})$  is differentiable, calculus provides the direction of steepest ascent.
- The Gradient:** The vector of partial derivatives with respect to all weights:

$$\nabla_{\mathbf{w}} J = \left[ \frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_d} \right]^T$$

# Optimization via Calculus

- Assuming  $J(\mathbf{w})$  is differentiable, calculus provides the direction of steepest ascent.
- The Gradient:** The vector of partial derivatives with respect to all weights:

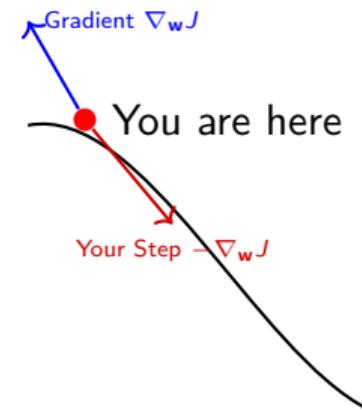
$$\nabla_{\mathbf{w}} J = \left[ \frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_d} \right]^T$$

- The negative gradient,  $-\nabla_{\mathbf{w}} J$ , points in the direction of steepest *descent*.

# Optimization via Calculus

- Assuming  $J(\mathbf{w})$  is differentiable, calculus provides the direction of steepest ascent.
- The Gradient:** The vector of partial derivatives with respect to all weights:

$$\nabla_{\mathbf{w}} J = \left[ \frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_d} \right]^T$$



- The negative gradient,  $-\nabla_{\mathbf{w}} J$ , points in the direction of steepest *descent*.
- We move opposite to the direction of gradient to get to the minimum of the loss function. This is called **Gradient Descent**

# The Supervised Learning Setup

**A typical supervised machine learning system consists of five key components:**

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .

# The Supervised Learning Setup

**A typical supervised machine learning system consists of five key components:**

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .

# The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .
- ③ **Parameters:** The unknown variables (e.g., weights  $\mathbf{w}$ ) within the model that determine its behavior. These must be learned.

# The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .
- ③ **Parameters:** The unknown variables (e.g., weights  $\mathbf{w}$ ) within the model that determine its behavior. These must be learned.
- ④ **Objective / Loss Function:** A metric  $J(\mathbf{w})$  (e.g., Squared Error) that quantifies the difference between the model's prediction  $\hat{y}$  and the true label  $y$ .

# The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .
- ③ **Parameters:** The unknown variables (e.g., weights  $\mathbf{w}$ ) within the model that determine its behavior. These must be learned.
- ④ **Objective / Loss Function:** A metric  $J(\mathbf{w})$  (e.g., Squared Error) that quantifies the difference between the model's prediction  $\hat{y}$  and the true label  $y$ .
- ⑤ **Learning Algorithm:** The optimization method (e.g., Gradient Descent) used to adjust the parameters  $\mathbf{w}$  to minimize the Loss Function.

# Recap: The Perceptron

## Formal Definition

A Perceptron is a linear classifier that maps a real-valued input vector  $\mathbf{x}$  to a binary output  $y$ .

$$y = \phi(\mathbf{w} \cdot \mathbf{x} + b)$$

Where  $\phi(z)$  is the **Step Function**:

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

# Recap: The Perceptron

## Formal Definition

A Perceptron is a linear classifier that maps a real-valued input vector  $\mathbf{x}$  to a binary output  $y$ .

$$y = \phi(\mathbf{w} \cdot \mathbf{x} + b)$$

Where  $\phi(z)$  is the **Step Function**:

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## Two Critical Limitations

### 1. Capacity

It can only solve linearly separable problems.

# Recap: The Perceptron

## Formal Definition

A Perceptron is a linear classifier that maps a real-valued input vector  $\mathbf{x}$  to a binary output  $y$ .

$$y = \phi(\mathbf{w} \cdot \mathbf{x} + b)$$

Where  $\phi(z)$  is the **Step Function**:

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## Two Critical Limitations

### 1. Capacity

It can only solve linearly separable problems.

We addressed this by stacking Perceptrons into a **Neural Network (MLP)**.

# Recap: The Perceptron

## Formal Definition

A Perceptron is a linear classifier that maps a real-valued input vector  $\mathbf{x}$  to a binary output  $y$ .

$$y = \phi(\mathbf{w} \cdot \mathbf{x} + b)$$

Where  $\phi(z)$  is the **Step Function**:

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## Two Critical Limitations

### 1. Capacity

It can only solve linearly separable problems.

We addressed this by stacking Perceptrons into a **Neural Network (MLP)**.

### 2. Threshold

The Step function ( $\phi$ ) has a "Harsh Threshold." Its derivative is either 0 or undefined.

# Recap: The Perceptron

## Formal Definition

A Perceptron is a linear classifier that maps a real-valued input vector  $\mathbf{x}$  to a binary output  $y$ .

$$y = \phi(\mathbf{w} \cdot \mathbf{x} + b)$$

Where  $\phi(z)$  is the **Step Function**:

$$\phi(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## Two Critical Limitations

### 1. Capacity

It can only solve linearly separable problems.

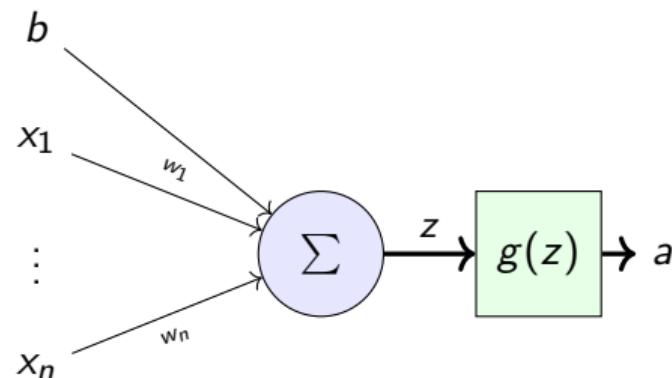
We addressed this by stacking Perceptrons into a **Neural Network (MLP)**.

### 2. Threshold

The Step function ( $\phi$ ) has a "Harsh Threshold." Its derivative is either 0 or undefined.

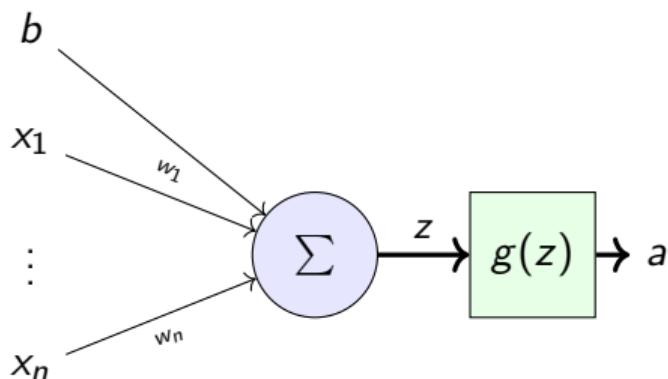
*How do we get a Gradient if the derivative is 0?*

# The Modern Neuron



**Figure:** A neuron with differentiable activation function

# The Modern Neuron

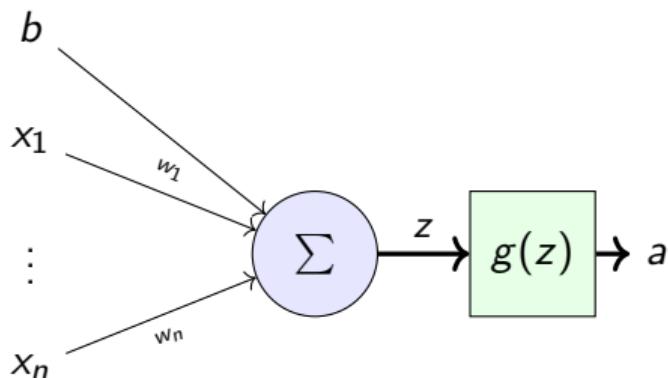


## Step 1: Linear Combination

- Calculates a weighted sum of inputs plus a bias.
- Output is  $z = \mathbf{w} \cdot \mathbf{x} + b$ .

Figure: A neuron with differentiable activation function

# The Modern Neuron



**Figure:** A neuron with differentiable activation function

## Step 1: Linear Combination

- Calculates a weighted sum of inputs plus a bias.
- Output is  $z = \mathbf{w} \cdot \mathbf{x} + b$ .

## Step 2: Differentiable Activation

- Passes the linear output  $z$  through an differentiable activation function  $a = g(z)$ .

# The Modern Neuron

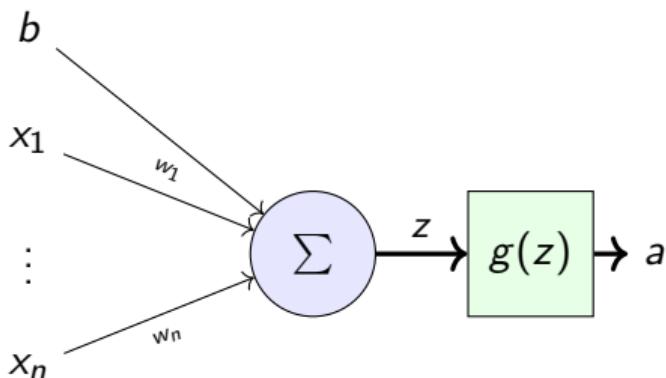


Figure: A neuron with differentiable activation function

## Step 1: Linear Combination

- Calculates a weighted sum of inputs plus a bias.
- Output is  $z = \mathbf{w} \cdot \mathbf{x} + b$ .

## Step 2: Differentiable Activation

- Passes the linear output  $z$  through an differentiable activation function  $a = g(z)$ .

## Why Differentiability is the Superpower?

Differentiable  $\rightarrow$  **gradient**. This gradient is used to adjust weights to fix errors. It is the core requirement for **backpropagation**.

# The Sigmoid Neuron

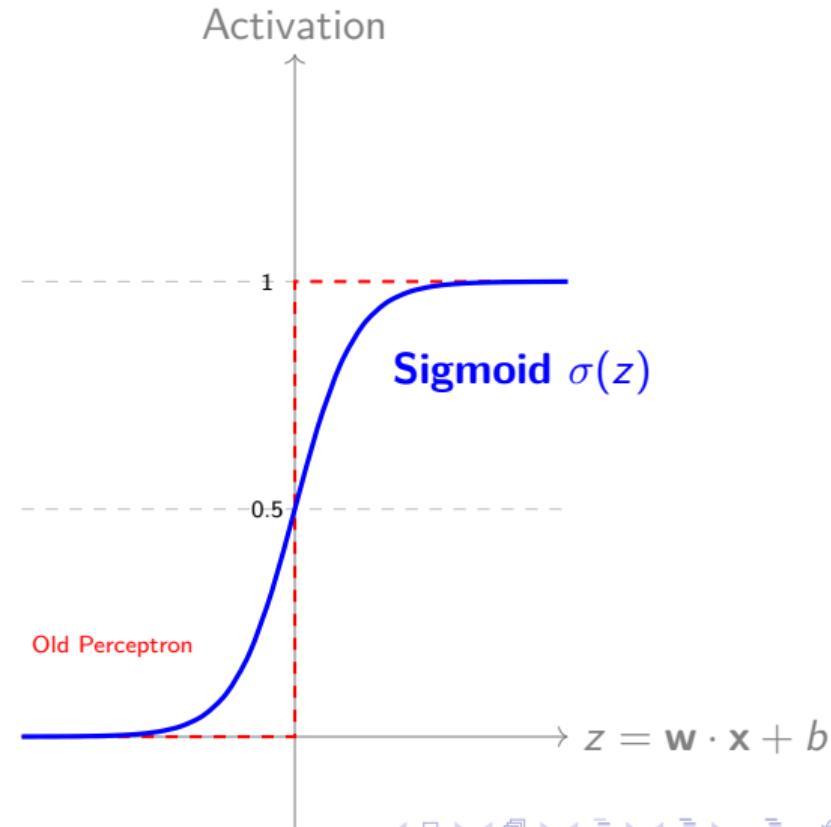
## The Definition

We replace the sharp Step Function with the smooth **Sigmoid Function** ( $\sigma$ ):

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

## Key Properties:

- **Continuous Output:** Returns a value between 0 and 1 (e.g., 0.73).
- **Probabilistic Interpretation:** Can be seen as  $P(y = 1|x)$ .
- **Differentiable:** The slope exists everywhere.
  - $\sigma'(z) \neq 0$  (mostly).



# Evolution of the Neuron: Side-by-Side Comparison

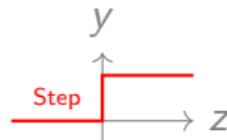
## 1. The Perceptron

**Equation:**

$$y = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b \geq 0 \\ 0 & \text{if } \mathbf{w} \cdot \mathbf{x} + b < 0 \end{cases}$$

**Characteristics:**

- **Output:** Binary {0, 1}.
- **Nature:** Hard Threshold.
- **Derivative:** 0 or Undefined.



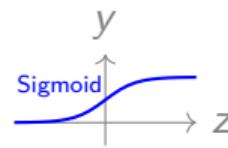
## 2. The Sigmoid Neuron

**Equation:**

$$y = \sigma(z) = \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}}$$

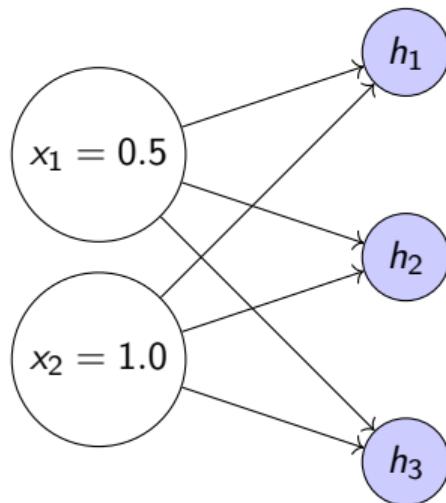
**Characteristics:**

- **Output:** Real value [0, 1].
- **Nature:** Smooth (S-Curve).
- **Derivative:** Non-zero ( $y(1 - y)$ ).



Let's compute the output of a hidden layer with 3 neurons, receiving input from 2 neurons, using **Sigmoid Activation**.

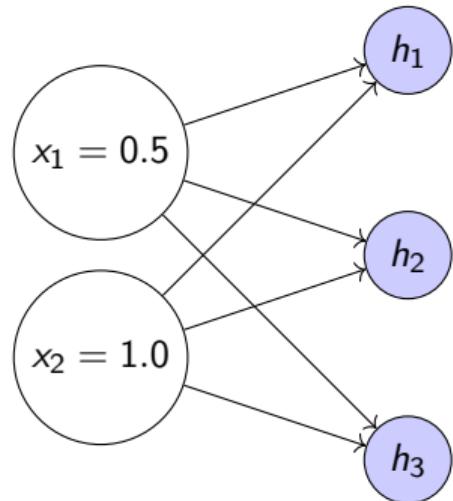
Let's compute the output of a hidden layer with 3 neurons, receiving input from 2 neurons, using **Sigmoid Activation**.



**Figure:** Input vector  $\mathbf{x}$  has 2 features.  
Hidden layer has 3 neurons.

Let's compute the output of a hidden layer with 3 neurons, receiving input from 2 neurons, using **Sigmoid Activation**.

### 1. Define inputs, weights, and biases:



$$\mathbf{x} = \begin{bmatrix} 0.5 \\ 1.0 \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} 0.2 & 0.7 \\ -0.4 & 0.1 \\ 0.9 & -0.3 \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} 0.1 \\ 0.2 \\ -0.5 \end{bmatrix}$$

Figure: Input vector  $\mathbf{x}$  has 2 features.  
Hidden layer has 3 neurons.

Let's compute the output of a hidden layer with 3 neurons, receiving input from 2 neurons, using **Sigmoid Activation**.

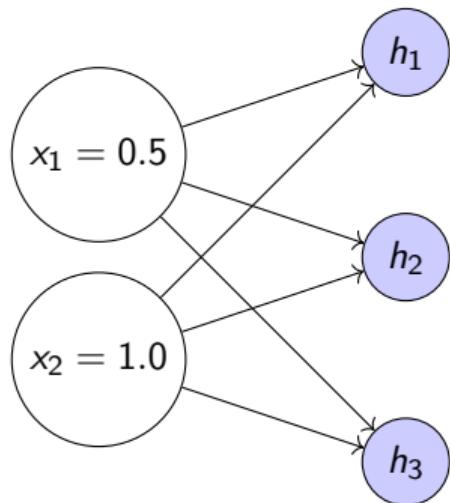


Figure: Input vector  $\mathbf{x}$  has 2 features.  
Hidden layer has 3 neurons.

**1. Define inputs, weights, and biases:**

$$\mathbf{x} = \begin{bmatrix} 0.5 \\ 1.0 \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} 0.2 & 0.7 \\ -0.4 & 0.1 \\ 0.9 & -0.3 \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} 0.1 \\ 0.2 \\ -0.5 \end{bmatrix}$$

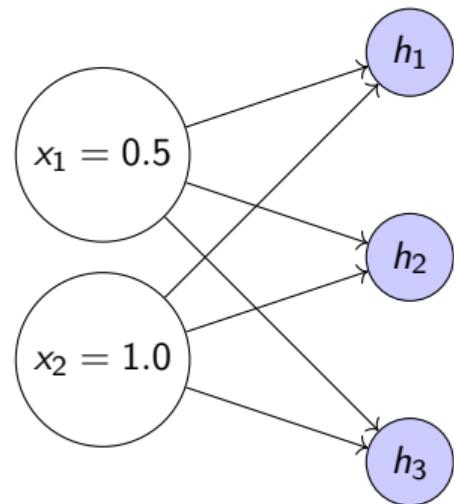
**2. Calculate the weighted sum  $\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$ :**

$$\mathbf{z} = \begin{bmatrix} 0.2 & 0.7 \\ -0.4 & 0.1 \\ 0.9 & -0.3 \end{bmatrix} \begin{bmatrix} 0.5 \\ 1.0 \end{bmatrix} + \begin{bmatrix} 0.1 \\ 0.2 \\ -0.5 \end{bmatrix} = \begin{bmatrix} 0.9 \\ 0.1 \\ -0.35 \end{bmatrix}$$

Let's compute the output of a hidden layer with 3 neurons, receiving input from 2 neurons, using **Sigmoid Activation**.

The activation function is Sigmoid:

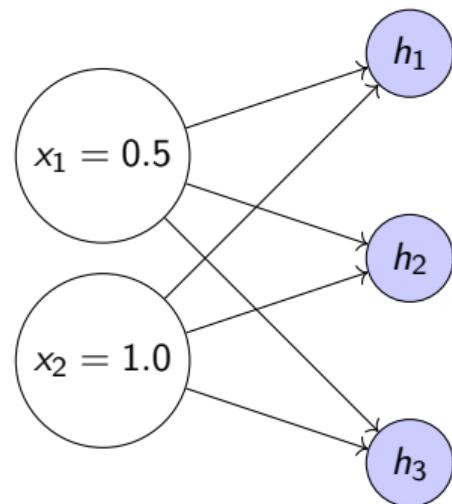
$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$



**Figure:** Input vector  $\mathbf{x}$  has 2 features.

Hidden layer has 3 neurons.

Let's compute the output of a hidden layer with 3 neurons, receiving input from 2 neurons, using **Sigmoid Activation**.



**Figure:** Input vector  $\mathbf{x}$  has 2 features.  
Hidden layer has 3 neurons.

The activation function is Sigmoid:

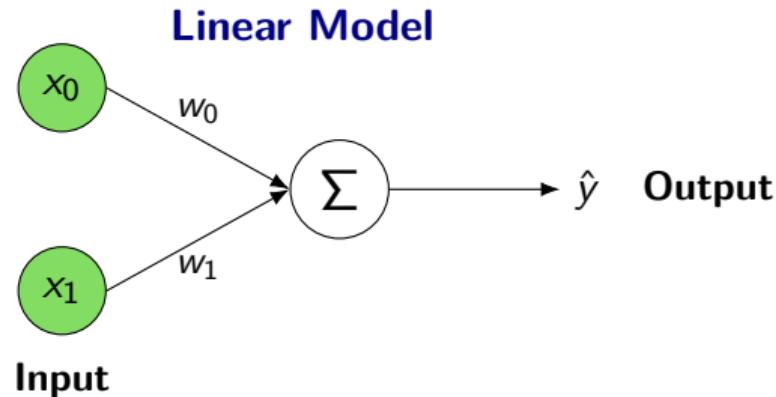
$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

**3. Apply Sigmoid activation  $\mathbf{a} = \sigma(\mathbf{z})$ :**

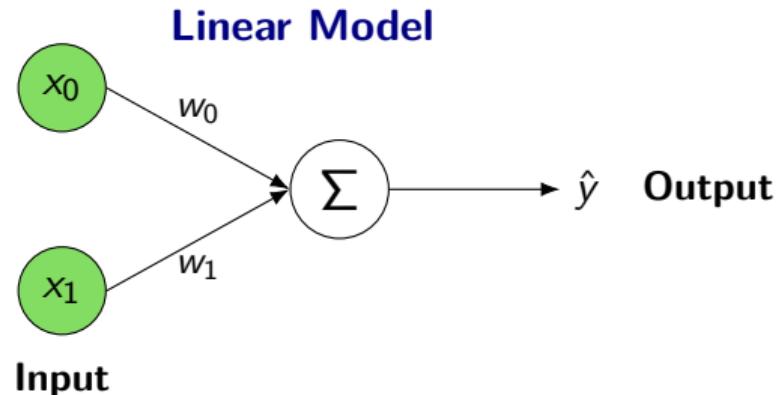
$$\mathbf{a} = \begin{bmatrix} \frac{1}{1+e^{-0.9}} \\ \frac{1}{1+e^{-0.1}} \\ \frac{1}{1+e^{0.35}} \end{bmatrix} \approx \begin{bmatrix} 0.71 \\ 0.52 \\ 0.41 \end{bmatrix}$$

This vector  $\mathbf{a}$  is the output of this layer.

# Linear Model vs Artificial Neural Networks

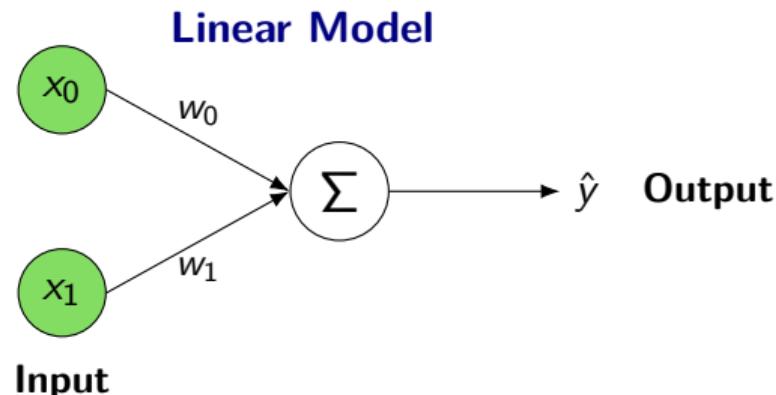


# Linear Model vs Artificial Neural Networks



$$\hat{y} = w_0x_0 + w_1x_1 = \sum w_i x_i$$

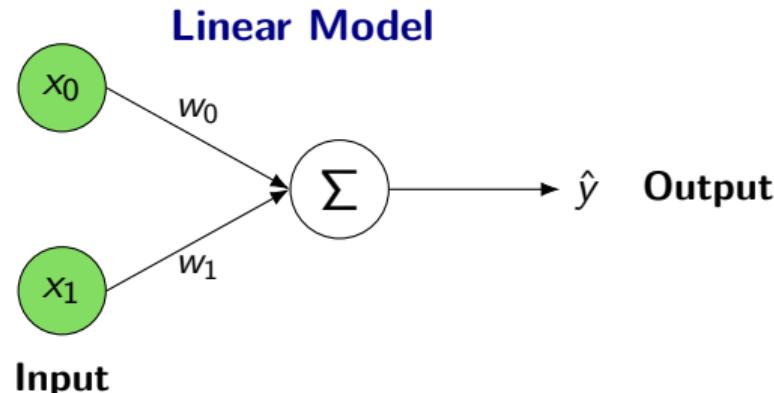
# Linear Model vs Artificial Neural Networks



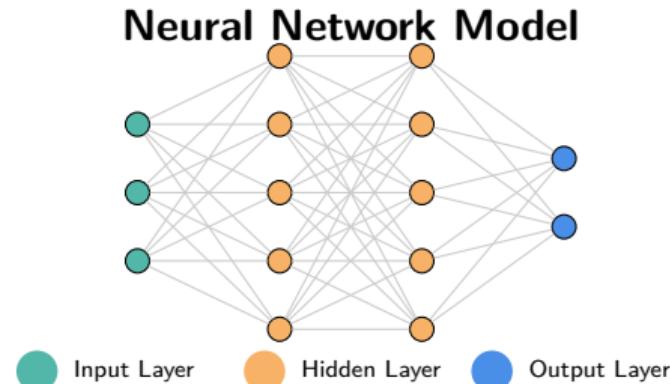
$$\hat{y} = w_0x_0 + w_1x_1 = \sum w_i x_i$$

A Linear Model is weighted sum of input features.

# Linear Model vs Artificial Neural Networks

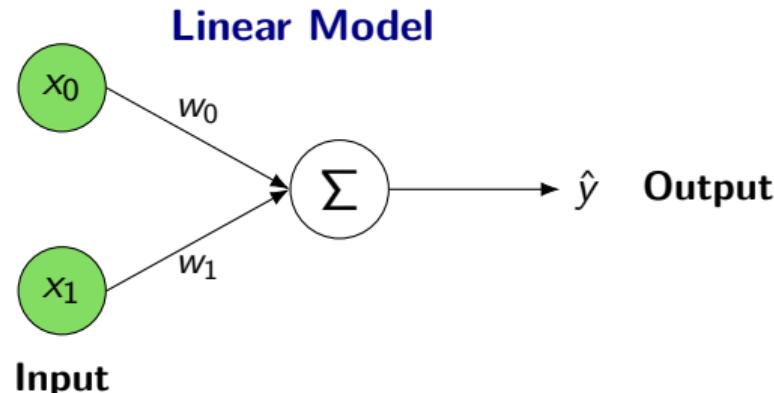


$$\hat{y} = w_0x_0 + w_1x_1 = \sum w_i x_i$$

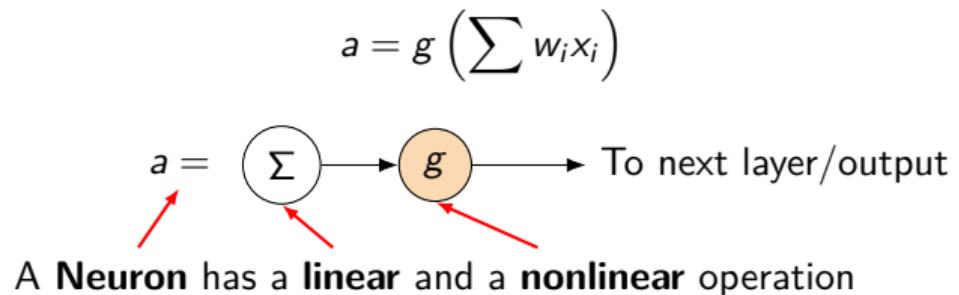
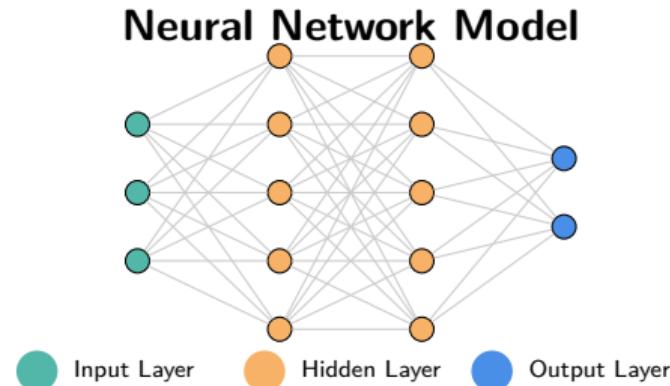


A Linear Model is weighted sum of input features.

# Linear Model vs Artificial Neural Networks



A Linear Model is weighted sum of input features.



# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .
- This is just another linear function:  $\mathbf{W}'\mathbf{x} + \mathbf{b}'$ .

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .
- This is just another linear function:  $\mathbf{W}'\mathbf{x} + \mathbf{b}'$ .

**A 100-layer linear network has the same power as a 1-layer network.**

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .
- This is just another linear function:  $\mathbf{W}'\mathbf{x} + \mathbf{b}'$ .

### With Non-Linearity ( $g$ ):

A 100-layer linear network has the same power as a 1-layer network.

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .
- This is just another linear function:  $\mathbf{W}'\mathbf{x} + \mathbf{b}'$ .

### With Non-Linearity ( $g$ ):

- The equation becomes:  
$$g(\mathbf{W}_2g(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2).$$

A 100-layer linear network has the same power as a 1-layer network.

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .
- This is just another linear function:  $\mathbf{W}'\mathbf{x} + \mathbf{b}'$ .

### With Non-Linearity ( $g$ ):

- The equation becomes:  
$$g(\mathbf{W}_2g(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2).$$
- This function **cannot** be simplified.

A 100-layer linear network has the same power as a 1-layer network.

# Why Do We Need Activation Functions?

## The Problem: Stacking Linear Layers is Useless

Without a non-linear activation function, a deep network simply collapses into a single linear model, no matter how many layers it has.

### Without Non-Linearity:

- A layer is a linear operation:  $\mathbf{Wx} + \mathbf{b}$ .
- Stacking them:  
$$L_2(L_1(\mathbf{x})) = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2.$$
- This simplifies to:  $(\mathbf{W}_2\mathbf{W}_1)\mathbf{x} + (\mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2)$ .
- This is just another linear function:  $\mathbf{W}'\mathbf{x} + \mathbf{b}'$ .

**A 100-layer linear network has the same power as a 1-layer network.**

### With Non-Linearity ( $g$ ):

- The equation becomes:  
$$g(\mathbf{W}_2g(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2).$$
- This function **cannot** be simplified.
- This allows the network to learn arbitrarily complex, "wiggly" functions.

# Representation Power of Multi-Layer Perceptron

# The Universal Approximation Theorem

## Theorem (Cybenko 1989, Hornik 1991)

Let  $\sigma(\cdot)$  be a non-constant, bounded, and continuous activation function (e.g., Sigmoid, Tanh, ReLU).

Then, for any continuous function  $f(x)$  defined on a compact subset of  $\mathbb{R}^n$  and for any error tolerance  $\epsilon > 0$ , there exists a Neural Network with **one hidden layer** containing a finite number of neurons that can approximate  $f(x)$  such that:

$$|F(x) - f(x)| < \epsilon \quad \forall x$$

# The Universal Approximation Theorem

## Theorem (Cybenko 1989, Hornik 1991)

Let  $\sigma(\cdot)$  be a non-constant, bounded, and continuous activation function (e.g., Sigmoid, Tanh, ReLU).

Then, for any continuous function  $f(x)$  defined on a compact subset of  $\mathbb{R}^n$  and for any error tolerance  $\epsilon > 0$ , there exists a Neural Network with **one hidden layer** containing a finite number of neurons that can approximate  $f(x)$  such that:

$$|F(x) - f(x)| < \epsilon \quad \forall x$$

### Implication:

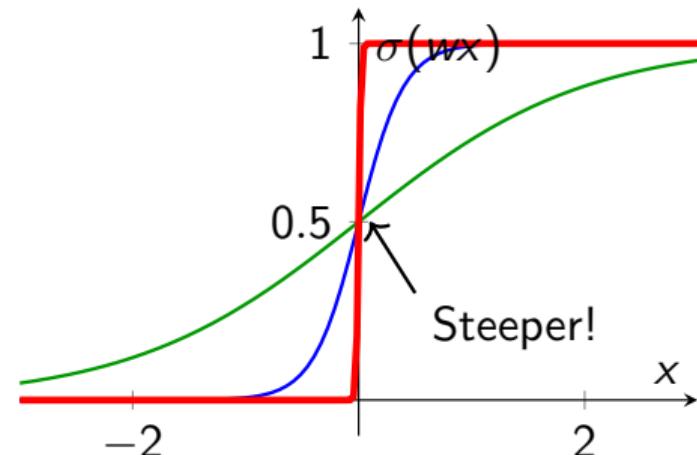
- Neural Networks are *universal function approximators*.
- In theory, a simple 2-layer network (1 hidden layer) can solve *any* problem, given enough neurons.

# Controlling Steepness with Weight ( $w$ )

## The Mechanism:

- Consider the sigmoid:  $\sigma(w \cdot x)$ .
- The weight  $w$  acts as a "gain" factor.
- **Small  $w$ :** The function is lazy and linear near the origin.
- **Large  $w$ :** The function transitions rapidly from 0 to 1.
- As  $w \rightarrow \infty$ , the sigmoid converges to a hard step.

## Effect of Weight magnitude

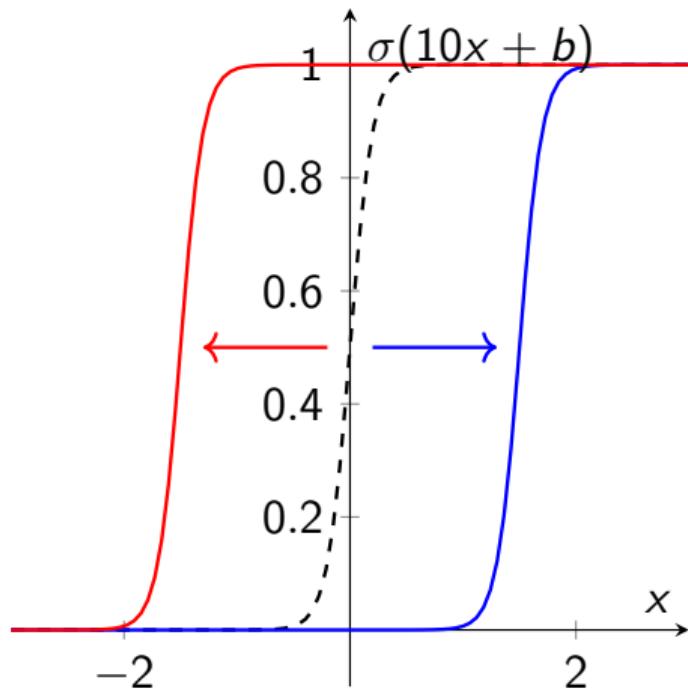


# Controlling Position with Bias ( $b$ )

Shifting the Step ( $w = 10$ )

## The Mechanism:

- Now consider:  $\sigma(w \cdot x + b)$ .
- The "step" occurs when the input to the sigmoid is 0.
- equation:  $wx + b = 0 \implies x = -b/w$ .
- Interpretation:** The center of the step is shifted to position  $s = -b/w$ .
- This allows us to place the "switch" anywhere on the x-axis.



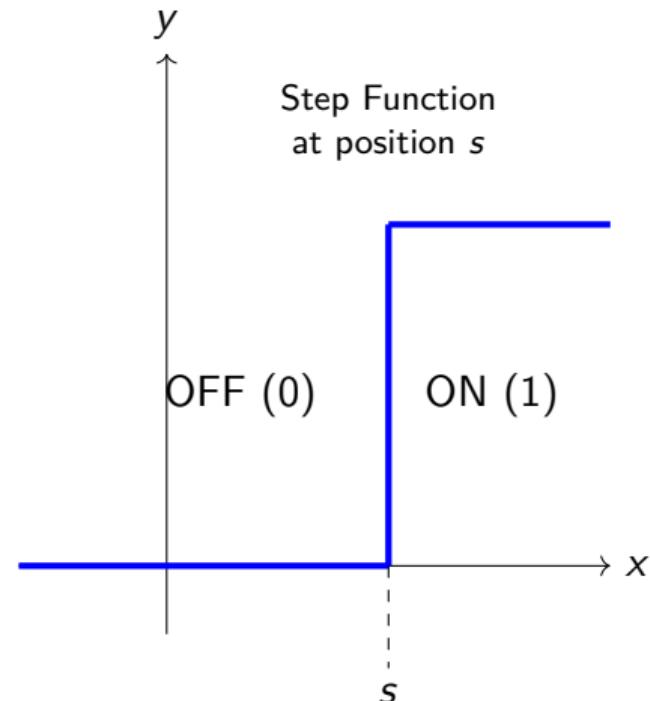
# Creating the "Hard" Step

## The Limit Definition:

- By combining a very large weight ( $w \rightarrow \infty$ ) and a specific bias, we can simulate a Heaviside Step Function  $H(x)$ .
- We define the step position as  $s$ .
- We set weights such that  $b = -w \cdot s$ .

$$\lim_{w \rightarrow \infty} \sigma(w(x - s)) = \begin{cases} 0 & \text{if } x < s \\ 1 & \text{if } x > s \end{cases}$$

**Why this matters:** This creates a "switch" that turns ON at exactly position  $x = s$ . We will use pairs of these switches to build "bumps."



# Constructing a "Bump"

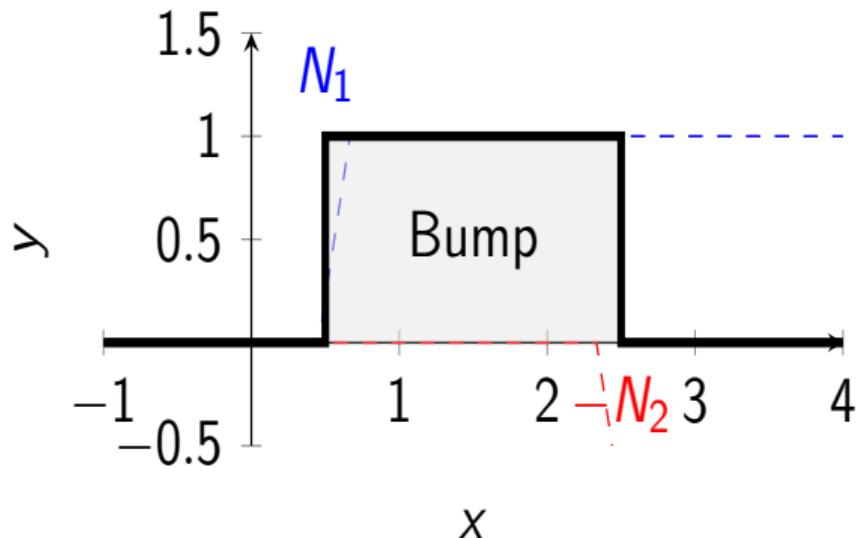
## The Tower Construction:

- ① Take Neuron 1 with step at  $s_1$ .
- ② Take Neuron 2 with step at  $s_2$ .
- ③ Subtract Neuron 2 from Neuron 1.

$$h(x) = \sigma(w(x - s_1)) - \sigma(w(x - s_2))$$

Result: A rectangular function (a "bump") that is non-zero only between  $s_1$  and  $s_2$ .

## Creating a Bump

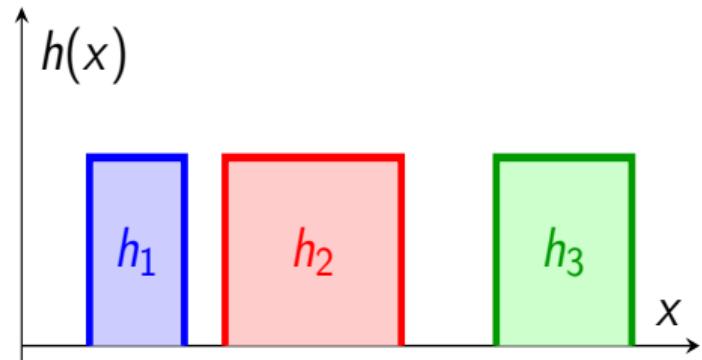


# The Building Blocks (Bumps)

## From Step to Bump:

- Recall: Two neurons create one "bump"  $h_j(x)$ .
- We can create  $N$  such bumps, scattered across the input space.
- **Key Idea:** Each bump is local. It is zero everywhere except for a specific region.
- We can independently control:
  - **Position:** Where the bump sits (via biases  $b$ ).
  - **Width:** How wide the bump is (via weights  $w$ ).

Independent "Hidden" Units



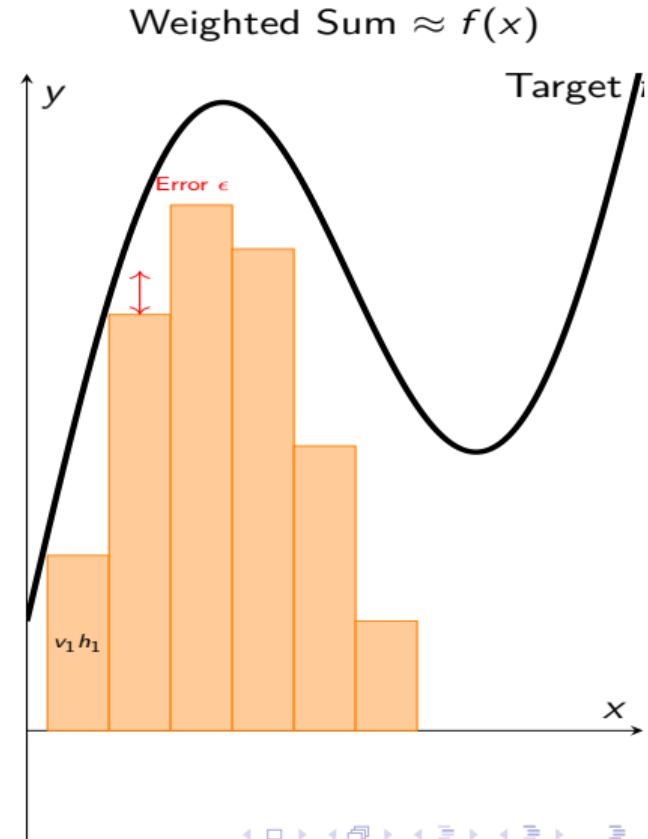
# Scaling and Summing to Fit

## Fitting the Function:

- We now have a set of bumps  $h_j(x)$ .
- The final output neuron computes a weighted sum:

$$F(x) = \sum_{j=1}^N v_j \cdot h_j(x)$$

- **Role of Output Weights ( $v_j$ ):** They scale the height of each bump to match the target function  $f(x)$  at that location.
- **Result:** A "histogram-like" approximation (Riemann Sum) of the curve.
- As  $N \rightarrow \infty$  (more bumps), the approximation error  $\rightarrow 0$ .



# Caveats and Practical Reality

## Theory vs. Practice

The Universal Approximation Theorem is an **existence theorem**, not a constructive one.

- **Existence:** It says a network *exists*. It does not tell us how to find the weights.
- **Efficiency:** The theorem requires a potentially **infinite** number of neurons in the hidden layer as the function becomes more complex.
- **Optimization:** Finding the optimal parameters using Gradient Descent is non-trivial (local minima, saddle points).
- **Why Deep Learning?** Instead of one massive wide layer, we use *deep* layers (many layers). This allows us to represent complex functions more efficiently (with fewer total parameters).

# Looking Ahead: From Existence to Discovery

## Module 2: Optimization & Learning

In the next section, we will learn how to efficiently navigate the loss landscape to find these parameters using calculus and adaptive algorithms:

- **Backpropagation:** The engine for computing gradients.
- **Gradient Descent:** The fundamental update rule.
- **Advanced Optimizers:** Adam, AdaDelta, and RMSProp.