

# DA6401:Introduction to Deep Learning

## Module 2A - BackPropagation in MLPs

Ganapathy Krishnamurthi

Department of Data Science and Artificial Intelligence , IIT Madras

# Story so far - The Supervised Learning Setup

**A typical supervised machine learning system consists of five key components:**

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .

# Story so far - The Supervised Learning Setup

**A typical supervised machine learning system consists of five key components:**

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .  
Typically comes from benchmarks and gold standard datasets.

# Story so far - The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .  
Typically comes from benchmarks and gold standard datasets.
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .

# Story so far - The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .  
Typically comes from benchmarks and gold standard datasets.
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .  
In our case, a Neural Network (i.e., a Multi-Layer Perceptron) with some fixed hyperparameters(i.e., number of layers, number of neurons, etc)

# Story so far - The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .  
Typically comes from benchmarks and gold standard datasets.
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .  
In our case, a Neural Network (i.e., a Multi-Layer Perceptron) with some fixed hyperparameters(i.e., number of layers, number of neurons, etc)
- ③ **Parameters:** The unknown variables within the model that determine its behavior. These must be learned.

# Story so far - The Supervised Learning Setup

A typical supervised machine learning system consists of five key components:

- ① **Data:** A collection of  $n$  labeled examples, denoted as  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n$ .  
Typically comes from benchmarks and gold standard datasets.
- ② **Model:** An approximation function that maps input  $x$  to output  $y$ .  
In our case, a Neural Network (i.e., a Multi-Layer Perceptron) with some fixed hyperparameters(i.e., number of layers, number of neurons, etc)
- ③ **Parameters:** The unknown variables within the model that determine its behavior. These must be learned.  
The weights and biases of the above defined MLP.

# Key Components to explore

- ④ **Objective / Loss Function:** A metric  $J(\mathbf{w})$  that quantifies the difference between the model's prediction  $\hat{y}$  and the true label  $y$ .

# Key Components to explore


- ④ **Objective / Loss Function:** A metric  $J(\mathbf{w})$  that quantifies the difference between the model's prediction  $\hat{y}$  and the true label  $y$ .
- ⑤ **Learning Algorithm:** The optimization method used to adjust the parameters  $\mathbf{w}$  to minimize the Loss Function.

# Output Activations & Loss Functions

# Recap: Regression vs. Classification

## 1. Regression


- **Goal:** Predict continuous quantities.
- **Example:** Predicting House Price & Tax.



# Recap: Regression vs. Classification

## 1. Regression


- **Goal:** Predict continuous quantities.
- **Example:** Predicting House Price & Tax.



*Regression: Continuous Output*

## 2. Classification

- **Goal:** Predict class probabilities.
- **Example:** Is the image a Cat or Dog?



*Classification: Probabilities*

# Regression: Activations and Loss

## Output Activation: Linear (Identity)

- $a(z) = z$  (Passes raw value through).

## Loss Function 1: Mean Squared Error (MSE)


Penalizes large errors heavily.

$$\mathcal{L}_{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

## Loss Function 2: Mean Absolute Error (MAE)

More robust to outliers.

$$\mathcal{L}_{MAE} = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$



# Classification: Multi-Class vs. Multi-Label

## Scenario 1: Multi-Class

- **Constraint:** Mutually Exclusive (One-Hot).
- *Example:* Digit is '3'. It cannot be '5'.

# Classification: Multi-Class vs. Multi-Label

## Scenario 1: Multi-Class

- **Constraint:** Mutually Exclusive (One-Hot).
- *Example:* Digit is '3'. It cannot be '5'.

## Scenario 2: Multi-Label

- **Constraint:** Independent / Non-exclusive.
- *Example:* Movie is 'Action' AND 'Sci-Fi'.


## Classification: Multi-Class vs. Multi-Label

## Scenario 1: Multi-Class

- **Constraint:** Mutually Exclusive (One-Hot).
  - *Example:* Digit is '3'. It cannot be '5'.

## Scenario 2: Multi-Label

- **Constraint:** Independent / Non-exclusive.
  - *Example:* Movie is 'Action' AND 'Sci-Fi'.




# Classification: Multi-Class vs. Multi-Label

## Scenario 1: Multi-Class

- **Constraint:** Mutually Exclusive (One-Hot).
- *Example:* Digit is '3'. It cannot be '5'.

## Scenario 2: Multi-Label


- **Constraint:** Independent / Non-exclusive.
- *Example:* Movie is 'Action' AND 'Sci-Fi'.



**Question:** In which scenario should the output probabilities sum to 1?

# Solutions: Activations

## Multi-Class (Mutually Exclusive)




*Softmax: Sums to 1*

### Activation: Softmax

$$\sigma(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

## Multi-Label (Independent)




*Sigmoid: Independent*

### Activation: Sigmoid

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

# Solutions: Loss Function

## Multi-Class (Mutually Exclusive)




Loss: Categorical Cross-Entropy

$$\mathcal{L} = - \sum_{c=1}^C y_c \log(\hat{y}_c)$$

(where  $y_c = 1$  for the true class)

## Multi-Label (Independent)



Loss: Binary Cross-Entropy

$$\mathcal{L} = - \frac{1}{M} \sum_{j=1}^M [y_j \log(\hat{y}_j) + (1 - y_j) \log(1 - \hat{y}_j)]$$

(Summed over all output nodes)

# Introduction to Backpropagation

# What is Backpropagation?

- **Definition:** Backpropagation ("backward propagation of errors") is an algorithm for computing the gradient of the loss function with respect to the weights.

# What is Backpropagation?

- **Definition:** Backpropagation ("backward propagation of errors") is an algorithm for computing the gradient of the loss function with respect to the weights.
- **Core Mechanism:** It applies the **Chain Rule** of calculus recursively from the output layer back to the input layer.

# What is Backpropagation?

- **Definition:** Backpropagation ("backward propagation of errors") is an algorithm for computing the gradient of the loss function with respect to the weights.
- **Core Mechanism:** It applies the **Chain Rule** of calculus recursively from the output layer back to the input layer.
- **Purpose:** It tells us how much each weight contributed to the error, allowing us to adjust them to reduce error.

## Use of Gradients

## The Goal: Minimize Loss $\mathcal{L}$

# Use of Gradients

## The Goal: Minimize Loss $\mathcal{L}$

- We view the Loss Surface as a landscape.

# Use of Gradients

## The Goal: Minimize Loss $\mathcal{L}$

- We view the Loss Surface as a landscape.
- The **Gradient** ( $\nabla_{\theta}\mathcal{L}$ ) points in the direction of steepest ascent (up the hill).


# Use of Gradients

## The Goal: Minimize Loss $\mathcal{L}$

- We view the Loss Surface as a landscape.
- The **Gradient** ( $\nabla_{\theta}\mathcal{L}$ ) points in the direction of steepest ascent (up the hill).
- To minimize error, we want to move in the direction of **Steepest Descent**.

$$\Delta W \propto -\frac{\partial \mathcal{L}}{\partial W}$$

*Backprop provides this gradient vector.*



# Objective of Backpropagation

Mathematically, for every weight  $w_{jk}^{(l)}$  and bias  $b_k^{(l)}$  in the network, we need to compute:

$$\frac{\partial \mathcal{L}}{\partial w_{jk}^{(l)}} \quad \text{and} \quad \frac{\partial \mathcal{L}}{\partial b_k^{(l)}}$$

Where:

- $\mathcal{L}$ : The Loss function (e.g., MSE, Cross-Entropy).
- $w_{jk}^{(l)}$ : Weight connecting neuron  $j$  in layer  $(l - 1)$  to neuron  $k$  in layer  $l$ .
- $b_k^{(l)}$ : Bias associated with neuron  $k$  in layer  $l$ .

# The Setup: Binary Classification


We want to train an MLP to predict the presence or absence of "Elephant" from an input image.

## Definitions:

- $y$ : True Label ( $1 = \text{Elephant}$ ,  $0 = \text{Not}$ ).
- $\hat{y}$ : Prediction ( $0.0 \rightarrow 1.0$ ).
- $z$ : Weighted sum into the output node.

## Chain Rule:

$$\frac{\partial \mathcal{L}}{\partial w_{jk}} = \underbrace{\frac{\partial \mathcal{L}}{\partial \hat{y}}}_{\text{Step 1}} \cdot \underbrace{\frac{\partial \hat{y}}{\partial z}}_{\text{Step 2}} \cdot \underbrace{\frac{\partial z}{\partial w_{jk}}}_{\text{Step 3}}$$



# Step 1: Derivative of Loss (BCE)

## Loss Function: Binary Cross-Entropy


$$\mathcal{L} = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

**The Derivative:** How does error change as prediction  $\hat{y}$  changes?

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = -\frac{y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}}$$

If we simplify this fraction:

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})}$$



## Step 2: Derivative of Activation (Sigmoid)


**Activation Function: Sigmoid** We squash the sum  $z$  to a probability:

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

**The Derivative:** The slope of the sigmoid curve at  $z$ .

$$\frac{\partial \hat{y}}{\partial z} = \sigma'(z) = \hat{y}(1 - \hat{y})$$

*Note: This matches the denominator from Step 1!*



# The "Magic" Simplification

Before Step 3, let's combine Step 1 and Step 2 to find the local error  $\delta$ .

$$\delta = \text{Step 1} \times \text{Step 2}$$

$$\delta = \left( \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot (\hat{y}(1 - \hat{y}))$$

$$\delta = \hat{y} - y$$

**Result:** The local error is just the difference between prediction and truth!

## Step 3: Derivative of Weights

**Question:** How much of the error is due to weight  $w_{jk}$ ?


Recall forward pass:

$$z = \dots + w_{jk} \cdot a_j + \dots$$

**Derivative:**

$$\frac{\partial z}{\partial w_{jk}} = a_j$$

Where  $a_j$  is the activation of the hidden neuron detecting "Elephant Features" (e.g., trunk shape).




# Final Gradient Calculation

Combining all parts for the weight update:

$$\frac{\partial \mathcal{L}}{\partial w_{jk}} = \underbrace{(\hat{y} - y)}_{\delta} \cdot \underbrace{a_j}_{\text{Input}}$$

## Interpretation:

- If Prediction  $\approx$  Truth, gradient is 0 (Stop learning).
- If Input  $a_j$  is high (feature present) AND error is high, update weight significantly.




# Extending to Deeper Layers: The Challenge

## What we've learned so far:

- We know how to compute gradients for weights connecting the **hidden layer** to the **output layer**:  $w_{jk}$
- Error at output:  $\delta_k = \hat{y} - y$

**The Question:** How do we update weights  $v_{ij}$  connecting the **input layer** to the **hidden layer**?

*Hint: We need to propagate the error backward through the network!*



# Step 1: Error at the Output Layer (Review)


**Output Layer Error:** We already computed this!

$$\delta_k = \frac{\partial \mathcal{L}}{\partial z_k} = \hat{y} - y$$

where  $z_k$  is the weighted sum into the output neuron.

*This is our starting point for backpropagation.*

**Next:** Propagate this error to the hidden layer.



## Step 2: Error at the Hidden Layer


Chain Rule to Hidden Neuron  $j$ :

$$\delta_j = \frac{\partial \mathcal{L}}{\partial z_j} = \frac{\partial \mathcal{L}}{\partial z_k} \cdot \frac{\partial z_k}{\partial a_j} \cdot \frac{\partial a_j}{\partial z_j}$$

Breaking it down:

- $\frac{\partial \mathcal{L}}{\partial z_k} = \delta_k$  (from output)
- $\frac{\partial z_k}{\partial a_j} = w_{jk}$  (weight connecting  $j$  to  $k$ )
- $\frac{\partial a_j}{\partial z_j} = \sigma'(z_j)$  (activation derivative)

$$\delta_j = \delta_k \cdot w_{jk} \cdot \sigma'(z_j)$$



Propagate error:  
 $\delta_j = \delta_k \cdot w_{jk} \cdot \sigma'(z_j)$

Error propagates through the weight  $w_{jk}$

## Step 3: Updating Input-to-Hidden Weights

**Gradient for weight  $v_{ij}$ :**

Now that we have  $\delta_j$  for the hidden layer, we can compute:


$$\frac{\partial \mathcal{L}}{\partial v_{ij}} = \delta_j \cdot x_i$$

where:

- $\delta_j$  = error at hidden neuron  $j$
- $x_i$  = input from pixel/feature  $i$

**Pattern Recognition:** This has the exact same form as before!

$$\frac{\partial \mathcal{L}}{\partial \text{weight}} = \text{error} \times \text{input}$$



Update  $v_{ij}$ :  
$$\frac{\partial \mathcal{L}}{\partial v_{ij}} = \delta_j \cdot x_i$$

# Vectorization: Why the Outer Product?

**From Scalar to Matrix:** In a fully connected layer, we don't update weights one by one. We update the entire weight matrix  $W^{(\ell)}$  at once.


**Individual Weight Gradient:**

$$\frac{\partial \mathcal{L}}{\partial w_{ij}} = \delta_j \cdot x_i$$

**Matrix Form:** To get the gradient for the whole matrix, we take the **Outer Product** of the error vector  $\delta$  and the input vector  $x^T$ :

$$\nabla_W \mathcal{L} = \delta x^T$$

**Visualization of  $\delta x^T$**



# Vectorization: Why the Outer Product?


## Individual Weight Gradient:

$$\frac{\partial \mathcal{L}}{\partial w_{ij}} = \delta_j \cdot x_i$$

**Matrix Form:** To get the gradient for the whole matrix, we take the **Outer Product** of the error vector  $\delta$  and the input vector  $x^T$ :

$$\nabla_W \mathcal{L} = \delta x^T$$

## Visualization of $\delta x^T$



**The Derivation:** Since  $z_i = \sum_j w_{ij} a_j$ , the partial derivative  $\frac{\partial z_i}{\partial w_{ij}}$  is only non-zero ( $a_j$ ) when the indices match. This maps every combination of error  $i$  and input  $j$  into its corresponding slot in the weight matrix.

# The General Backpropagation Recipe

For a network with  $L$  layers, we backpropagate errors from layer  $L$  to layer 1:

## 1. Output Layer (Layer $L$ ):

$$\delta^{(L)} = \hat{y} - y$$

## 2. Hidden Layers ( $\ell = L - 1, L - 2, \dots, 2$ ):

$$\delta^{(\ell)} = (W^{(\ell+1)})^T \delta^{(\ell+1)} \odot \sigma'(z^{(\ell)})$$

(where  $\odot$  is element-wise multiplication)

## 3. Weight Updates (all layers):

$$\frac{\partial \mathcal{L}}{\partial W^{(\ell)}} = \delta^{(\ell)} (a^{(\ell-1)})^T$$

*The error "flows backward" through the weights, scaled by activation derivatives.*

# Key Insights: Why Backpropagation Works

## ① Chain Rule is Everything:

We decompose the complex derivative into simple, local computations at each layer.

## ② Reuse Computations:

Each  $\delta$  is computed once and reused for all weights in that layer. This makes backprop efficient!

## ③ Error Attribution:

Weights that contributed more to the error ( $\text{high } \delta \cdot \text{input}$ ) get larger updates.

## ④ Vanishing/Exploding Gradients:

As we go deeper,  $\delta$  can become very small (vanishing) or very large (exploding) due to repeated multiplication by  $w \cdot \sigma'(z)$ . This is why deep networks can be hard to train!


# Automatic Differentiation in Frameworks

# The Mechanics of Autograd

To compute gradients, PyTorch builds a **Dynamic Directed Acyclic Graph (DAG)**.

- **Tensors (Nodes)**: Store the data and the accumulated gradient (`.grad`).
- **Function Objects (Edges)**: Every operation (add, mul, exp) creates a `grad_fn` that knows its own derivative.
- **Leaf Tensors**: Tensors created by the user (weights  $w$ , bias  $b$ , input  $x$ ).


Rule: Gradient at Node = Incoming Gradient  $\times$  Local Derivative



# The Anatomy of a Graph Node

In PyTorch, every operation (e.g.,  $+$ ,  $\times$ ,  $\exp$ ) is a node that implements a `forward()` and a `backward()` method.

**FORWARD: Compute *out***




**BACKWARD:**  $\frac{\partial L}{\partial a} = \frac{\partial L}{\partial out} \cdot \frac{\partial out}{\partial a}$

## A 5-Step Example: $L = \sigma(w \cdot x + b)$

Let's decompose the operation  $L = \frac{1}{1+e^{-(wx+b)}}$  into atomic steps:

- |   |                       |           |
|---|-----------------------|-----------|
| ① | $z_1 = w \times x$    | Mul       |
| ② | $z_2 = z_1 + b$       | Add       |
| ③ | $z_3 = -z_2$          | Neg       |
| ④ | $z_4 = \exp(z_3)$     | Exp       |
| ⑤ | $L = \frac{1}{1+z_4}$ | Inv + Add |




# Step-by-Step Backpropagation Trace

Starting with  $\frac{\partial L}{\partial L} = 1$ , we move backwards node-by-node:

Node	Local Grad ( $\frac{\partial \text{out}}{\partial \text{in}}$ )	Chain Rule Expression	Result
$L = \frac{1}{1+z_4}$	$\frac{\partial L}{\partial z_4} = -\frac{1}{(1+z_4)^2}$	$\frac{\partial L}{\partial L} \cdot \frac{\partial L}{\partial z_4} = 1 \cdot (-L^2)$	$\frac{\partial L}{\partial z_4} = -L^2$
$z_4 = \exp(z_3)$	$\frac{\partial z_4}{\partial z_3} = \exp(z_3)$	$\frac{\partial L}{\partial z_4} \cdot \frac{\partial z_4}{\partial z_3} = (-L^2) \cdot z_4$	$\frac{\partial L}{\partial z_3} = -L^2 z_4$
$z_3 = -z_2$	$\frac{\partial z_3}{\partial z_2} = -1$	$\frac{\partial L}{\partial z_3} \cdot \frac{\partial z_3}{\partial z_2} = (-L^2 z_4) \cdot (-1)$	$\frac{\partial L}{\partial z_2} = L^2 z_4$
$z_2 = z_1 + b$	$\frac{\partial z_2}{\partial b} = 1$	$\frac{\partial L}{\partial z_2} \cdot \frac{\partial z_2}{\partial b} = (L^2 z_4) \cdot 1$	$\frac{\partial L}{\partial b} = L^2 z_4$
$z_1 = w \cdot x$	$\frac{\partial z_1}{\partial w} = x$	$\frac{\partial L}{\partial z_1} \cdot \frac{\partial z_1}{\partial w} = (L^2 z_4) \cdot x$	$\frac{\partial L}{\partial w} = L^2 z_4 x$


**Key Point:** The weight gradient  $\frac{\partial L}{\partial w}$  is the product of all local derivatives encountered on the path from  $L$  back to  $w$ .



# Why the Graph? Gradient Accumulation

If a variable  $w$  affects  $L$  through **two different paths**,  
the gradients are **summed** at the junction.


- Path 1 gives  $\frac{\partial L}{\partial p_1}$
- Path 2 gives  $\frac{\partial L}{\partial p_2}$
- **Total:**  $\frac{\partial L}{\partial w} = \frac{\partial L}{\partial p_1} + \frac{\partial L}{\partial p_2}$



# Step-by-Step Example: $L = (w \cdot x + b - y)^2$

Let's trace a forward pass with 5 operations:

- ①  $z_1 = w \cdot x$  (Multiplication)
- ②  $z_2 = z_1 + b$  (Addition)
- ③  $z_3 = z_2 - y$  (Subtraction)
- ④  $L = z_3^2$  (Power)



# The Backward Pass: Message Passing

When we call `L.backward()`, the gradient  $g = \frac{\partial L}{\partial L} = 1$  flows backward:

- **At Power Node ( $z_3^2$ ):** Local  $\frac{\partial L}{\partial z_3} = 2z_3$ . Grad sent back:  $1 \cdot 2z_3$ .
- **At Subtraction Node ( $z_2 - y$ ):** Local  $\frac{\partial z_3}{\partial z_2} = 1$ . Grad sent back:  $(2z_3) \cdot 1$ .
- **At Addition Node ( $z_1 + b$ ):** Local  $\frac{\partial z_2}{\partial z_1} = 1$  and  $\frac{\partial z_2}{\partial b} = 1$ .
- **At Multiplication Node ( $w \cdot x$ ):** Local  $\frac{\partial z_1}{\partial w} = x$ . Final Grad:  $(2z_3 \cdot 1 \cdot 1) \cdot x$ .

# Key Features of the Graph Approach

## Efficiency:

- Only local derivatives are calculated at each node.
- Intermediate results ( $z_1, z_2, z_3$ ) are cached during forward pass.

## Gradient Accumulation:

- If a tensor is used in multiple operations, the gradients from different paths are **summed** at that node.
- This is how PyTorch handles branching.

# The Forward and Backward Pass

## 1. Forward Pass:

- Data flows from left to right.
- Each operation saves **intermediate values** (tensors) needed for the gradient.
- Each tensor points to a `grad_fn` (the recipe for its derivative).

# The Forward and Backward Pass

## 1. Forward Pass:

- Data flows from left to right.
- Each operation saves **intermediate values** (tensors) needed for the gradient.
- Each tensor points to a `grad_fn` (the recipe for its derivative).

## 2. Backward Pass (`loss.backward()`):

- The graph is traversed **in reverse**.
- Local gradients are computed using the saved values.
- Gradients are multiplied by the incoming "gradient from above" (Chain Rule).

# From Math to Code

## The Graph Logic:

- ① Define tensors with `requires_grad=True`.
- ② Perform operations (PyTorch builds the graph).
- ③ Call `.backward()`.
- ④ Access `.grad` attribute.

```
import torch

x = torch.tensor(1.0, requires_grad=True)
w = torch.tensor(2.0, requires_grad=True)
b = torch.tensor(0.5, requires_grad=True)

# Forward pass (Building Graph)
z = w * x + b
loss = z**2

# Backward pass (Traversing Graph)
loss.backward()

print(w.grad) # dLoss/dw
```

# Illustration of Backpropagation: Setup

## Tiny Network for Hand Calculation:


- 2 input features:  $x_1 = 0.5, x_2 = 1.0$
- 2 hidden neurons
- 1 output neuron (binary classification)
- Sigmoid activation:  $\sigma(z) = \frac{1}{1+e^{-z}}$
- True label:  $y = 1$

## Initial Weights:

$$v_{11} = 0.5, \quad v_{12} = 0.5$$

$$v_{21} = 0.5, \quad v_{22} = 0.5$$

$$w_1 = 1.0, \quad w_2 = 1.0$$



# Step 1: Forward Pass - Hidden Layer

**Compute Hidden Layer Activations:**

**Hidden Neuron 1:**

$$z_1 = v_{11}x_1 + v_{21}x_2 = 0.5 \times 0.5 + 0.5 \times 1.0$$

$$z_1 = 0.25 + 0.5 = \boxed{0.75}$$


$$h_1 = \sigma(z_1) = \frac{1}{1 + e^{-0.75}} \approx \boxed{0.679}$$

**Hidden Neuron 2:**

$$z_2 = v_{12}x_1 + v_{22}x_2 = 0.5 \times 0.5 + 0.5 \times 1.0$$

$$z_2 = 0.25 + 0.5 = \boxed{0.75}$$

$$h_2 = \sigma(z_2) = \frac{1}{1 + e^{-0.75}} \approx \boxed{0.679}$$



Both hidden neurons have  
activation  $\approx 0.679$

## Step 2: Forward Pass - Output Layer

**Compute Output:**

**Weighted Sum:**

$$z_{\text{out}} = w_1 h_1 + w_2 h_2 = 1.0 \times 0.679 + 1.0 \times 0.679$$


$$z_{\text{out}} = 0.679 + 0.679 = \boxed{1.358}$$

**Final Prediction:**

$$\hat{y} = \sigma(z_{\text{out}}) = \frac{1}{1 + e^{-1.358}} \approx \boxed{0.796}$$

**Loss (Binary Cross-Entropy):**

$$\begin{aligned}\mathcal{L} &= -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})] \\ &= -[1 \times \log(0.796) + 0 \times \log(0.204)] \\ &= -\log(0.796) \approx \boxed{0.228}\end{aligned}$$



Prediction: 0.796

Error:  $0.796 - 1 = -0.204$

## Step 3: Backward Pass - Output Error

### Compute Output Layer Error:

Recall the simplified formula for sigmoid + BCE:


$$\delta_{\text{out}} = \hat{y} - y$$

### Calculation:

$$\delta_{\text{out}} = 0.796 - 1 = -0.204$$

Negative error means we predicted too low (should increase output).

**Next Step:** Use this to compute gradients for  $w_1$  and  $w_2$ .



Output error computed!  
Now backpropagate...

## Step 4: Gradients for Output Weights ( $w_1, w_2$ )


**Gradient Formula:**  $\frac{\partial \mathcal{L}}{\partial w_i} = \delta_{\text{out}} \cdot h_i$

For  $w_1$ :

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_1} &= \delta_{\text{out}} \cdot h_1 = (-0.204) \times 0.679 \\ &= \boxed{-0.139}\end{aligned}$$

For  $w_2$ :

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_2} &= \delta_{\text{out}} \cdot h_2 = (-0.204) \times 0.679 \\ &= \boxed{-0.139}\end{aligned}$$



Gradients computed!  
Both weights need to increase

Negative gradients mean we should increase both weights!

## Step 5: Hidden Layer Errors ( $\delta_1, \delta_2$ )

### Backpropagate Error to Hidden Layer:

Formula:  $\delta_j = \delta_{\text{out}} \cdot w_j \cdot \sigma'(z_j)$

First, compute  $\sigma'(z_j) = h_j(1 - h_j)$ :


$$\sigma'(z_1) = 0.679 \times (1 - 0.679) = 0.218$$

### For Hidden Neuron 1:

$$\begin{aligned}\delta_1 &= \delta_{\text{out}} \cdot w_1 \cdot \sigma'(z_1) \\ &= (-0.204) \times 1.0 \times 0.218 = \boxed{-0.044}\end{aligned}$$

### For Hidden Neuron 2:

$$\begin{aligned}\delta_2 &= \delta_{\text{out}} \cdot w_2 \cdot \sigma'(z_2) \\ &= (-0.204) \times 1.0 \times 0.218 = \boxed{-0.044}\end{aligned}$$



Hidden errors computed!  
Error propagated backward

## Step 6: Gradients for Hidden Weights ( $v_{ij}$ )

**Gradient Formula:**  $\frac{\partial \mathcal{L}}{\partial v_{ij}} = \delta_j \cdot x_i$

Weights to Hidden Neuron 1:

$$\frac{\partial \mathcal{L}}{\partial v_{11}} = \delta_1 \cdot x_1 = (-0.044) \times 0.5 = \boxed{-0.022}$$

Weights to Hidden Neuron 2:

$$\frac{\partial \mathcal{L}}{\partial v_{12}} = \delta_2 \cdot x_1 = (-0.044) \times 0.5 = \boxed{-0.022}$$

$$\frac{\partial \mathcal{L}}{\partial v_{21}} = \delta_1 \cdot x_2 = (-0.044) \times 1.0 = \boxed{-0.044}$$

$$\frac{\partial \mathcal{L}}{\partial v_{22}} = \delta_2 \cdot x_2 = (-0.044) \times 1.0 = \boxed{-0.044}$$

**All Gradients Computed!**

All hidden layer weights have negative gradients → should increase.

# Connecting Backpropagation to Optimization

## The Story So Far:

- **Forward Pass:** We mapped inputs to predictions and calculated the Loss  $\mathcal{L}$ .
- **Backpropagation:** We used the chain rule to find exactly how each weight contributes to that loss ( $\nabla W$ ).

**The Missing Link:** Now that we have the gradients, how do we actually **improve** the model?

- ① The gradient  $\frac{\partial \mathcal{L}}{\partial W}$  points in the direction of steepest *increase*.
- ② To minimize error, we must move in the **opposite direction**.

## Next Step: Gradient Descent Algorithm and its variants

We will now see how this update rule allows the model to "learn" by iteratively descending the loss landscape.