

# DA6401:Introduction to Deep Learning

## Module 1A-Artificial Neuron Models

Ganapathy Krishnamurthi

Department of Data Science and Artificial Intelligence , IIT Madras

## Course Details

- **Instructor:** Prof. Ganapathy Krishnamurthi
- **Email:**  
gankrish@dsai.iitm.ac.in
- **Venue:** CRC302 (In-person)
- **Slot C:**
  - Mon (10:00), Tue (09:00)
  - Wed (08:00), Fri (12:00)

## Communication

- **Channels:** Moodle
- **Dates:** Jan 20 – May 5, 2026

### Note on Lectures

Classes are primarily in-person. Online links will be provided if physical attendance is not feasible.

# Getting Ready for the Course

## Course Prerequisites

- **Foundational Courses:**

- DA5000: **Mathematical Foundations of Data Science**
- DA5400: **Foundations of Machine Learning**

# Getting Ready for the Course

## Course Prerequisites

- **Foundational Courses:**

- DA5000: [Mathematical Foundations of Data Science](#)
- DA5400: [Foundations of Machine Learning](#)

- **Technical Proficiency:**

- Strong command of **Python**.
- Assignments will rely heavily on [NumPy](#) and [PyTorch](#).

# Grading Scheme & Key Dates

# Grading Scheme & Key Dates

Component	Weight	Date
Quiz 1 (MCQ)	20%	<b>Feb 18, 2026</b>
Quiz 2 (MCQ)	20%	<b>Mar 25, 2026</b>

# Grading Scheme & Key Dates

Component	Weight	Date
Quiz 1 (MCQ)	20%	<b>Feb 18, 2026</b>
Quiz 2 (MCQ)	20%	<b>Mar 25, 2026</b>
Assignments ( $4 \times 5\%$ )	20%	<i>Throughout Sem</i>

# Grading Scheme & Key Dates

Component	Weight	Date
Quiz 1 (MCQ)	20%	<b>Feb 18, 2026</b>
Quiz 2 (MCQ)	20%	<b>Mar 25, 2026</b>
Assignments ( $4 \times 5\%$ )	20%	<i>Throughout Sem</i>
<b>Final Exam*</b>	<b>40%</b>	<b>May 8, 2026</b>

*\*The final exam will be handwritten.*



# Course Roadmap: Fundamentals (Modules 1-3)

- **Module 1: Introduction**

- History, Perceptrons, and the XOR problem.
- **Multilayer Perceptrons (MLPs)** and representation power.

# Course Roadmap: Fundamentals (Modules 1-3)

- **Module 1: Introduction**

- History, Perceptrons, and the XOR problem.
- **Multilayer Perceptrons (MLPs)** and representation power.

- **Module 2: Optimization**

- Backpropagation derivation.
- Optimizers: **SGD, Adam, RMSProp**.
- Initialization techniques (Xavier/He).

# Course Roadmap: Fundamentals (Modules 1-3)

- **Module 1: Introduction**

- History, Perceptrons, and the XOR problem.
- **Multilayer Perceptrons (MLPs)** and representation power.

- **Module 2: Optimization**

- Backpropagation derivation.
- Optimizers: **SGD, Adam, RMSProp**.
- Initialization techniques (Xavier/He).

- **Module 3: Training Methodology**

- Regularization: L2, Dropout, Early Stopping.
- Handling Class Imbalance.

# Course Roadmap: Specialized Architectures (Modules 4-6)

- **Module 4: Computer Vision (CNNs)**

- Convolutions, Pooling, LeNet to ResNet.
- Object Detection (YOLO) & Segmentation (U-Net).

# Course Roadmap: Specialized Architectures (Modules 4-6)

- **Module 4: Computer Vision (CNNs)**

- Convolutions, Pooling, LeNet to ResNet.
- Object Detection (YOLO) & Segmentation (U-Net).

- **Module 5: Sequence Models (NLP)**

- Word Embeddings (Word2Vec).
- RNNs, LSTMs, and GRUs.
- Introduction to Attention.

# Course Roadmap: Specialized Architectures (Modules 4-6)

- **Module 4: Computer Vision (CNNs)**

- Convolutions, Pooling, LeNet to ResNet.
- Object Detection (YOLO) & Segmentation (U-Net).

- **Module 5: Sequence Models (NLP)**

- Word Embeddings (Word2Vec).
- RNNs, LSTMs, and GRUs.
- Introduction to Attention.

- **Module 6: Generative AI**

- **Transformers** & Self-Attention.
- Generative Models: VAEs, GANs, and **Diffusion Models**.

## Textbooks

- **Deep Learning** – Ian Goodfellow, Yoshua Bengio, Aaron Courville (MIT Press).
- **Dive into Deep Learning** – Zhang et al. (d2l.ai).

*Questions?*

# Machine Learning Landscape

From Traditional Machine Learning to Deep Learning

---



# Machine Learning

## Arthur Samuel (1959)

Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed.

# Machine Learning

## Arthur Samuel (1959)

Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed.

**Learn directly from data - No explicitly programmed rules**

# Machine Learning

## Arthur Samuel (1959)

Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed.

**Learn directly from data - No explicitly programmed rules**

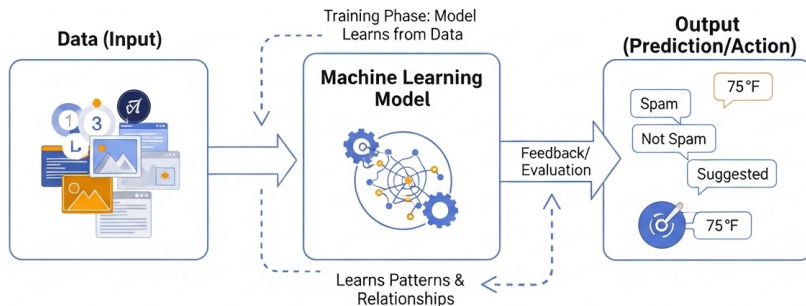


Figure: ML: Learning from Data (Source: ImageGen)

# Common Machine Learning Applications



Figure: Housing Price Prediction

# Common Machine Learning Applications



Figure: Housing Price Prediction




Figure: Autonomous vehicles(Lane following)

# Common Machine Learning Applications



Figure: Housing Price Prediction




Figure: Autonomous vehicles(Lane following)

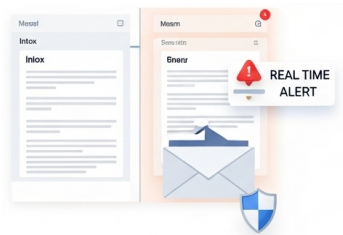


Figure: Message Sentiment Analysis

# Common Machine Learning Applications

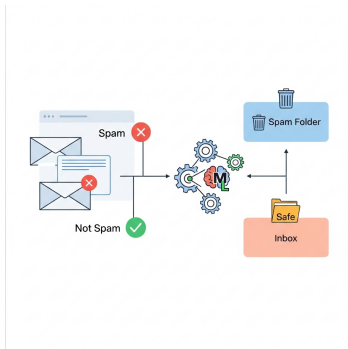


Figure: Spam Email Filtering

# Common Machine Learning Applications

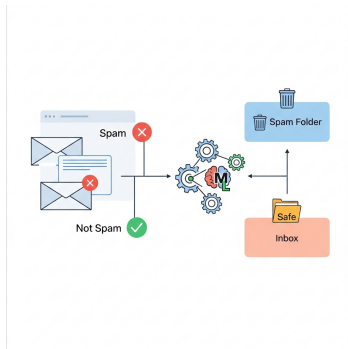


Figure: Spam Email Filtering




Figure: Recommendation engines



# Common Machine Learning Applications

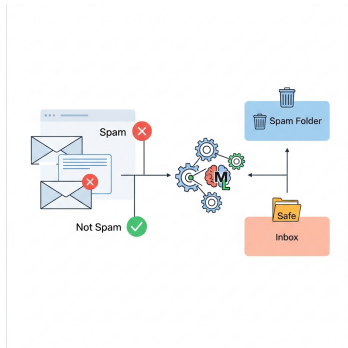


Figure: Spam Email Filtering




Figure: Recommendation engines




Figure: Image Recognition

## The Data & Process

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models



# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models

### Linear Models

- *Core Idea:* Assumes linear relationships.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models

### Linear Models

- *Core Idea:* Assumes linear relationships.
- *Examples:* Linear/Logistic Regression.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models

### Linear Models

- *Core Idea*: Assumes linear relationships.
- *Examples*: Linear/Logistic Regression.
- *Traits*: Fast, interpretable, but simple.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models

### Linear Models

- *Core Idea*: Assumes linear relationships.
- *Examples*: Linear/Logistic Regression.
- *Traits*: Fast, interpretable, but simple.

### Tree-Based Models

- *Core Idea*: Learns 'if-then-else' rules.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models

### Linear Models

- *Core Idea*: Assumes linear relationships.
- *Examples*: Linear/Logistic Regression.
- *Traits*: Fast, interpretable, but simple.

### Tree-Based Models

- *Core Idea*: Learns 'if-then-else' rules.
- *Examples*: Random Forest, Gradient Boosting.

# Machine Learning: Key Concepts

## The Data & Process

### Focus: Structured Data

- Data in tables (rows & columns).
- Predict a target from features.

### Method: Manual Feature Engineering

- Human experts design features.
- Model learns from these engineered features.
- Success depends entirely on feature quality.

## The Models

### Linear Models

- *Core Idea*: Assumes linear relationships.
- *Examples*: Linear/Logistic Regression.
- *Traits*: Fast, interpretable, but simple.

### Tree-Based Models

- *Core Idea*: Learns 'if-then-else' rules.
- *Examples*: Random Forest, Gradient Boosting.
- *Traits*: Captures non-linearity, powerful.

# Limitations of Traditional Machine Learning

# Limitations of Traditional Machine Learning




Figure: Feature Engineering Bottleneck  
(Source: ImageGen)

- **The Feature Engineering Bottleneck:** Time-consuming, requires domain expertise, and is often suboptimal.



# Limitations of Traditional Machine Learning




Figure: Feature Engineering Bottleneck  
(Source: ImageGen)

- **The Feature Engineering Bottleneck:** Time-consuming, requires domain expertise, and is often suboptimal.
- **Inability to handle high-dimensional data:** Struggles with images, audio, or raw text.

# Limitations of Traditional Machine Learning




Figure: Feature Engineering Bottleneck  
(Source: ImageGen)

- **The Feature Engineering Bottleneck:** Time-consuming, requires domain expertise, and is often suboptimal.
- **Inability to handle high-dimensional data:** Struggles with images, audio, or raw text.
- **Limited representation learning:** Learns shallow patterns, not deep, hierarchical features.

# Limitations of Traditional Machine Learning




Figure: Feature Engineering Bottleneck  
(Source: ImageGen)

- **The Feature Engineering Bottleneck:** Time-consuming, requires domain expertise, and is often suboptimal.
- **Inability to handle high-dimensional data:** Struggles with images, audio, or raw text.
- **Limited representation learning:** Learns shallow patterns, not deep, hierarchical features.
- **Performance plateaus:** Adding more data yields diminishing returns after a certain point.

# Deep Learning - A New Paradigm

From Traditional Machine Learning to Deep Learning

---

# The Shift to Deep Learning: A New Paradigm

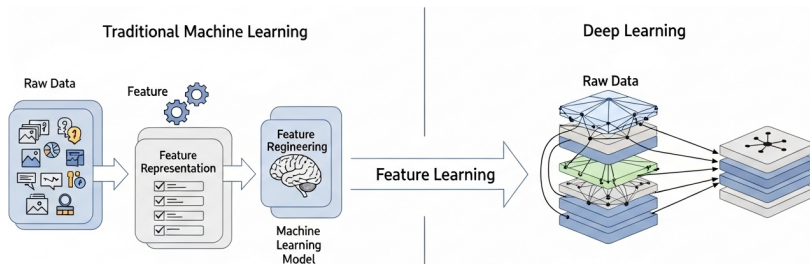


Figure: ML needs manual feature extraction; DL learns them automatically.

# The Shift to Deep Learning: A New Paradigm




Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.

# The Shift to Deep Learning: A New Paradigm




Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.
- Handles images, text, and audio.

# The Shift to Deep Learning: A New Paradigm




Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.
- Handles images, text, and audio.
- **No manual feature engineering.**



# The Shift to Deep Learning: A New Paradigm

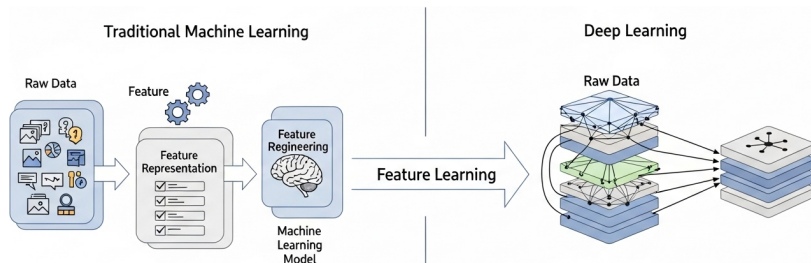


Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.
- Handles images, text, and audio.
- **No manual feature engineering.**

## Why this Paradigm Shift?

# The Shift to Deep Learning: A New Paradigm

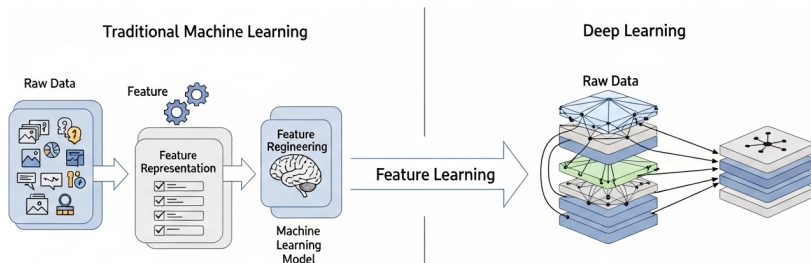


Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.
- Handles images, text, and audio.
- **No manual feature engineering.**

## Why this Paradigm Shift?

- Solves the feature engineering bottleneck.

# The Shift to Deep Learning: A New Paradigm




Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.
- Handles images, text, and audio.
- **No manual feature engineering.**

## Why this Paradigm Shift?

- Solves the feature engineering bottleneck.
- Automates representation learning.

# The Shift to Deep Learning: A New Paradigm

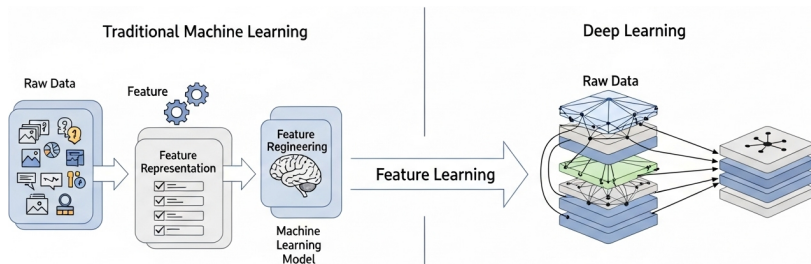


Figure: ML needs manual feature extraction; DL learns them automatically.

## Deep Learning: The Core Idea

- Models learn directly from raw data.
- Handles images, text, and audio.
- **No manual feature engineering.**

## Why this Paradigm Shift?

- Solves the feature engineering bottleneck.
- Automates representation learning.
- Unlocks new, complex capabilities.

# Popular Deep Learning Applications

# Popular Deep Learning Applications




Figure: Computer vision

# Popular Deep Learning Applications



Figure: Computer vision




Figure: Speech recognition and Audio synthesis

# Popular Deep Learning Applications




Figure: Computer vision



Figure: Natural language processing




Figure: Speech recognition and Audio synthesis



# Popular Deep Learning Applications




Figure: Computer vision



Figure: Natural language processing




Figure: Speech recognition and Audio synthesis




Figure: Game playing and robotics

**Machine Learning**

**Deep Learning**

## Machine Learning

- **Feature Extraction:**  
Requires **manual** engineering.

## Deep Learning

- **Feature Extraction:**  
Learns features **automatically**.

## Machine Learning

- **Feature Extraction:**  
Requires **manual** engineering.
- **Representation:**  
Learns **shallow** patterns.

## Deep Learning

- **Feature Extraction:**  
Learns features **automatically**.
- **Representation:**  
Builds **hierarchical** features.

## Machine Learning

- **Feature Extraction:**  
Requires **manual** engineering.
- **Representation:**  
Learns **shallow** patterns.
- **Data Scalability:**  
Works well on **small to medium** data.

## Deep Learning

- **Feature Extraction:**  
Learns features **automatically**.
- **Representation:**  
Builds **hierarchical** features.
- **Data Scalability:**  
Excels with **very large** datasets.

## Machine Learning

- **Feature Extraction:**  
Requires **manual** engineering.
- **Representation:**  
Learns **shallow** patterns.
- **Data Scalability:**  
Works well on **small to medium** data.

## Deep Learning

- **Feature Extraction:**  
Learns features **automatically**.
- **Representation:**  
Builds **hierarchical** features.
- **Data Scalability:**  
Excels with **very large** datasets.
- **Training Paradigm:**  
Enables **end-to-end** learning.

# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*

# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*





# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*



## Big Data

- Massive labeled datasets.
- The "fuel" for hungry models.
- e.g., ImageNet, Wikipedia.

# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*



## Big Data

- Massive labeled datasets.
- The "fuel" for hungry models.
- e.g., ImageNet, Wikipedia.



## Computational Power

# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*



## Big Data

- Massive labeled datasets.
- The "fuel" for hungry models.
- e.g., ImageNet, Wikipedia.



## Computational Power

- Rise of GPUs (parallel power).
- Made deep training feasible.
- Drastically cut training time.

# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*



## Big Data

- Massive labeled datasets.
- The "fuel" for hungry models.
- e.g., ImageNet, Wikipedia.



## Computational Power

- Rise of GPUs (parallel power).
- Made deep training feasible.
- Drastically cut training time.



## Better Algorithms

# What Enabled the Deep Learning?

*A convergence of three key factors created the "perfect environment"*



## Big Data

- Massive labeled datasets.
- The "fuel" for hungry models.
- e.g., ImageNet, Wikipedia.



## Computational Power


- Rise of GPUs (parallel power).
- Made deep training feasible.
- Drastically cut training time.



## Better Algorithms

- Architectural innovations.
- e.g., CNNs, Transformers.
- Smarter training methods.

# The AI Venn Diagram




# From Biology to Mathematics

Building Blocks of a Neural Network

---

# Inspiration from the Brain



**Figure:** A biological neuron (Source: Wikipedia)



# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.

# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.
- Integrates signals in the **Soma/Cell body**.

# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.
- Integrates signals in the **Soma/Cell body**.
- If threshold is met, it **"Fires"**.

# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.
- Integrates signals in the **Soma/Cell body**.
- If threshold is met, it **"Fires"**.
- Sends output signal via **Axon**.

# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.
- Integrates signals in the **Soma/Cell body**.
- If threshold is met, it **"Fires"**.
- Sends output signal via **Axon**.

### Important: Inspiration, Not a Replica

- Brain is vastly more complex.

# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.
- Integrates signals in the **Soma/Cell body**.
- If threshold is met, it **"Fires"**.
- Sends output signal via **Axon**.

## Important: Inspiration, Not a Replica

- Brain is vastly more complex.
- Biological learning is not backpropagation.

# Inspiration from the Brain




Figure: A biological neuron (Source: Wikipedia)

## A Simple Model of a Neuron:

- Receives signals via **Dendrites**.
- Integrates signals in the **Soma/Cell body**.
- If threshold is met, it **"Fires"**.
- Sends output signal via **Axon**.

## Important: Inspiration, Not a Replica

- Brain is vastly more complex.
- Biological learning is not backpropagation.
- This is a **useful abstraction**, not a perfect model.

# The McCulloch-Pitts Neuron (1943)



# The McCulloch-Pitts Neuron (1943)


## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."

# The McCulloch-Pitts Neuron (1943)

## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."



**Figure:** The M-P Neuron Model: Sum inputs and compare to a threshold  $\theta$ .


# The McCulloch-Pitts Neuron (1943)

## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."

### How it Works:

- Takes multiple binary inputs ( $x_i \in \{0, 1\}$ ).



**Figure:** The M-P Neuron Model: Sum inputs and compare to a threshold  $\theta$ .


# The McCulloch-Pitts Neuron (1943)

## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."

### How it Works:

- Takes multiple binary inputs ( $x_i \in \{0, 1\}$ ).
- Calculates their sum.



**Figure:** The M-P Neuron Model: Sum inputs and compare to a threshold  $\theta$ .


# The McCulloch-Pitts Neuron (1943)

## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."

### How it Works:

- Takes multiple binary inputs ( $x_i \in \{0, 1\}$ ).
- Calculates their sum.
- The output is 1 if the sum is greater than or equal to a fixed threshold ( $\theta$ ).



**Figure:** The M-P Neuron Model: Sum inputs and compare to a threshold  $\theta$ .


# The McCulloch-Pitts Neuron (1943)

## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."

### How it Works:

- Takes multiple binary inputs ( $x_i \in \{0, 1\}$ ).
- Calculates their sum.
- The output is 1 if the sum is greater than or equal to a fixed threshold ( $\theta$ ).
- The output is 0 otherwise.




**Figure:** The M-P Neuron Model: Sum inputs and compare to a threshold  $\theta$ .

# The McCulloch-Pitts Neuron (1943)

## The Core Idea: A Simple Biological Logic Gate

The first mathematical model of a neuron. It operates on a simple principle: if enough inputs are active, the neuron "fires."



**Figure:** The M-P Neuron Model: Sum inputs and compare to a threshold  $\theta$ .

### How it Works:

- Takes multiple binary inputs ( $x_i \in \{0, 1\}$ ).
- Calculates their sum.
- The output is 1 if the sum is greater than or equal to a fixed threshold ( $\theta$ ).
- The output is 0 otherwise.

### Significance:


- Proved that networks of these simple units could compute any logical function (e.g., AND, OR).

# M-P Neuron Examples: Logical Functions

- Inputs are binary:  $x \in \{0, 1\}$ .

## 1. The AND Function

*(Fires only if all inputs are 1)*






# M-P Neuron Examples: Logical Functions

- Inputs are binary:  $x \in \{0, 1\}$ .

## 1. The AND Function

*(Fires only if all inputs are 1)*




# M-P Neuron Examples: Logical Functions

- Inputs are binary:  $x \in \{0, 1\}$ .


## 1. The AND Function

*(Fires only if all inputs are 1)*



## 2. The OR Function

*(Fires if at least one input is 1)*




# M-P Neuron Examples: Logical Functions

- Inputs are binary:  $x \in \{0, 1\}$ .


## 1. The AND Function

*(Fires only if all inputs are 1)*



## 2. The OR Function

*(Fires if at least one input is 1)*




# M-P Neuron Examples: Logical Functions

- Inputs are binary:  $x \in \{0, 1\}$ .


## 1. The AND Function

*(Fires only if all inputs are 1)*



## 2. The OR Function

*(Fires if at least one input is 1)*



### Observation

By simply changing the **threshold** ( $\theta$ ), the same physical structure performs a completely different logical function.

# Real-World Analogy: The "Movie Night" Decision

## **The Scenario:**

You are deciding whether to go to a movie tonight ( $y$ ).

# Real-World Analogy: The "Movie Night" Decision

## The Scenario:

You are deciding whether to go to a movie tonight ( $y$ ).

## The Inputs (Factors):

$x_1$ : Is the movie a comedy? (1=Yes, 0=No)

$x_2$ : Is a friend coming? (1=Yes, 0=No)

$x_3$ : Is the ticket cheap? (1=Yes, 0=No)

# Real-World Analogy: The "Movie Night" Decision

## The Scenario:

You are deciding whether to go to a movie tonight ( $y$ ).

## The Inputs (Factors):

$x_1$ : Is the movie a comedy? (1=Yes, 0=No)

$x_2$ : Is a friend coming? (1=Yes, 0=No)

$x_3$ : Is the ticket cheap? (1=Yes, 0=No)

## The Personality (Threshold):

You are **easy-going**. You only need **2 out of 3** things to go right to be convinced.

$$\theta = 2$$

# Real-World Analogy: The "Movie Night" Decision

## The Scenario:

You are deciding whether to go to a movie tonight ( $y$ ).

## The Inputs (Factors):

$x_1$ : Is the movie a comedy? (1=Yes, 0=No)


$x_2$ : Is a friend coming? (1=Yes, 0=No)

$x_3$ : Is the ticket cheap? (1=Yes, 0=No)

## The Personality (Threshold):

You are **easy-going**. You only need **2 out of 3** things to go right to be convinced.

$$\theta = 2$$





# Real-World Analogy: The "Movie Night" Decision

## The Scenario:

You are deciding whether to go to a movie tonight ( $y$ ).

## The Inputs (Factors):

$x_1$ : Is the movie a comedy? (1=Yes, 0=No)


$x_2$ : Is a friend coming? (1=Yes, 0=No)

$x_3$ : Is the ticket cheap? (1=Yes, 0=No)

## The Personality (Threshold):

You are **easy-going**. You only need **2 out of 3** things to go right to be convinced.

$$\theta = 2$$



## Situation

It's a Horror movie ( $x_1 = 0$ ), but your friend is coming ( $x_2 = 1$ ) and it's cheap ( $x_3 = 1$ ).

# Real-World Analogy: The "Movie Night" Decision

## The Scenario:

You are deciding whether to go to a movie tonight ( $y$ ).

## The Inputs (Factors):

$x_1$ : Is the movie a comedy? (1=Yes, 0=No)


$x_2$ : Is a friend coming? (1=Yes, 0=No)

$x_3$ : Is the ticket cheap? (1=Yes, 0=No)

## The Personality (Threshold):

You are **easy-going**. You only need **2 out of 3** things to go right to be convinced.

$$\theta = 2$$



## Situation

It's a Horror movie ( $x_1 = 0$ ), but your friend is coming ( $x_2 = 1$ ) and it's cheap ( $x_3 = 1$ ).

$$\text{Sum} = 0 + 1 + 1 = 2$$

$$2 \geq 2 \implies \text{Fire (Go to Movie)}$$

# Geometric Interpretation: Drawing a Line

## The Mathematical Boundary

Recall the firing condition for 2 inputs:

$$x_1 + x_2 \geq \theta$$

# Geometric Interpretation: Drawing a Line

## The Mathematical Boundary

Recall the firing condition for 2 inputs:

$$x_1 + x_2 \geq \theta$$

This inequality defines a **half-plane**. The boundary is the line:

$$x_1 + x_2 = \theta \quad \Rightarrow \quad \mathbf{x}_2 = -\mathbf{x}_1 + \theta$$

# Geometric Interpretation: Drawing a Line

## The Mathematical Boundary


Recall the firing condition for 2 inputs:

$$x_1 + x_2 \geq \theta$$

This inequality defines a **half-plane**. The boundary is the line:

$$x_1 + x_2 = \theta \quad \Rightarrow \quad x_2 = -x_1 + \theta$$

- Points **above** the line: Neuron Fires (1).
- Points **below** the line: Neuron Silent (0).



# Geometric Interpretation: Drawing a Line

## The Mathematical Boundary


Recall the firing condition for 2 inputs:

$$x_1 + x_2 \geq \theta$$

This inequality defines a **half-plane**. The boundary is the line:

$$x_1 + x_2 = \theta \quad \Rightarrow \quad x_2 = -x_1 + \theta$$

- Points **above** the line: Neuron Fires (1).
- Points **below** the line: Neuron Silent (0).



## Key Takeaway: Linear Separability

The MP Neuron is a **Linear Classifier**. It works by drawing a *single straight line* to separate the "Yes" answers from the "No" answers.

# Where the Model Breaks: The "Used Car" Dilemma

## The Scenario:

You are buying a used car. You have a threshold of **2** "Yes" factors to say "I'll buy it."

# Where the Model Breaks: The "Used Car" Dilemma

## The Scenario:

You are buying a used car. You have a threshold of 2 "Yes" factors to say "I'll buy it."

## The Features ( $x$ ):

- 1 Is the Engine good? ( $x_1$ )
- 2 Are the Brakes working? ( $x_2$ )
- 3 Is there a Cup Holder? ( $x_3$ )



# Where the Model Breaks: The "Used Car" Dilemma


## The Scenario:

You are buying a used car. You have a threshold of **2** "Yes" factors to say "I'll buy it."

## The Features ( $x$ ):

- 1 Is the Engine good? ( $x_1$ )
- 2 Are the Brakes working? ( $x_2$ )
- 3 Is there a Cup Holder? ( $x_3$ )

## MP Neuron Logic ( $\theta = 2$ )



# Where the Model Breaks: The "Used Car" Dilemma


## The Scenario:

You are buying a used car. You have a threshold of **2** "Yes" factors to say "I'll buy it."

## The Features ( $x$ ):

- 1 Is the Engine good? ( $x_1$ )
- 2 Are the Brakes working? ( $x_2$ )
- 3 Is there a Cup Holder? ( $x_3$ )

## MP Neuron Logic ( $\theta = 2$ )



*In the real world, Brakes  $\gg$  Cup Holder.*

# Where the Model Breaks: The "Used Car" Dilemma


## The Scenario:

You are buying a used car. You have a threshold of 2 "Yes" factors to say "I'll buy it."

## The Features ( $x$ ):

- 1 Is the Engine good? ( $x_1$ )
- 2 Are the Brakes working? ( $x_2$ )
- 3 Is there a Cup Holder? ( $x_3$ )

## MP Neuron Logic ( $\theta = 2$ )



*In the real world, Brakes  $\gg$  Cup Holder.*

## Failure Mode

Consider a car with **No Brakes**, but a Good Engine and a Cup Holder:

- $x_1 = 1, x_2 = 0, x_3 = 1$
- $\text{Sum} = 1 + 0 + 1 = 2$

**Result:** The MP Neuron says "**BUY**" because it treats a *Cup Holder* as equal to *Brakes*.

# Critical Limitations of the M-P Neuron

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.
- *Missing piece: **Weights**.*

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.
- *Missing piece: **Weights**.*



### No Automatic Learning



# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.
- *Missing piece: **Weights**.*



### No Automatic Learning

- Threshold ( $\theta$ ) is fixed.
- Must be set manually.

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.
- *Missing piece: **Weights**.*



### No Automatic Learning

- Threshold ( $\theta$ ) is fixed.
- Must be set manually.
- *Missing piece: **A Learning Rule**.*

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.
- *Missing piece: **Weights**.*



### No Automatic Learning

- Threshold ( $\theta$ ) is fixed.
- Must be set manually.
- *Missing piece: **A Learning Rule**.*



### Restricted to Binary I/O

# Critical Limitations of the M-P Neuron

## A Calculator, Not a Learner

The M-P neuron was a fixed logic gate. It could compute, but it could not **learn**.



### No Feature Importance

- Equally important features.
- Can't prioritize key features.
- *Missing piece: **Weights**.*



### No Automatic Learning

- Threshold ( $\theta$ ) is fixed.
- Must be set manually.
- *Missing piece: **A Learning Rule**.*




### Restricted to Binary I/O

- Inputs must be 0 or 1.
- Outputs are only 0 or 1.
- Can't handle **real-valued data**.


# The Leap to Learning: The Perceptron (1958)

# The Leap to Learning: The Perceptron (1958)



**Figure:** The Perceptron: Introduces learnable weights ( $w_i$ ) on each input.

# The Leap to Learning: The Perceptron (1958)




**Figure:** The Perceptron: Introduces learnable weights ( $w_i$ ) on each input.

## Innovation 1: Weighted Inputs

- Each input is multiplied by a **weight**.
- Weights represent feature importance.
- The model can now **prioritize** signals.

# The Leap to Learning: The Perceptron (1958)



**Figure:** The Perceptron: Introduces learnable weights ( $w_i$ ) on each input.

## Innovation 1: Weighted Inputs


- Each input is multiplied by a **weight**.
- Weights represent feature importance.
- The model can now **prioritize** signals.

## Innovation 2: The Learning Rule

- Compares prediction ( $\hat{y}$ ) to true label ( $y$ ).
- If wrong, it adjusts the weights to correct the error.



# The Leap to Learning: The Perceptron (1958)



**Figure:** The Perceptron: Introduces learnable weights ( $w_i$ ) on each input.

## Innovation 1: Weighted Inputs

- Each input is multiplied by a **weight**.
- Weights represent feature importance.
- The model can now **prioritize** signals.

## Innovation 2: The Learning Rule

- Compares prediction ( $\hat{y}$ ) to true label ( $y$ ).
- If wrong, it adjusts the weights to correct the error.

## Significance: The Birth of Learning Machines

For the first time, **a machine could automatically learn** to classify patterns, laying the foundation for all of modern deep learning.

## McCulloch-Pitts (The Old Way)

- **Inputs:** Binary  $x_i \in \{0, 1\}$
- **Weights:** None (all 1)
- **Threshold:** Fixed  $\theta$

$$y = 1 \quad \text{if} \quad \sum_{i=1}^n x_i \geq \theta$$

# Mathematical Formulation: From Fixed to Flexible

## McCulloch-Pitts (The Old Way)

- **Inputs:** Binary  $x_i \in \{0, 1\}$
- **Weights:** None (all 1)
- **Threshold:** Fixed  $\theta$

$$y = 1 \quad \text{if} \quad \sum_{i=1}^n x_i \geq \theta$$

## The Perceptron (The New Way)

- **Inputs:** Real numbers  $x_i \in \mathbb{R}$
- **Weights:** Learnable  $w_i \in \mathbb{R}$
- **Bias:** Learnable  $b$  (or  $w_0$ )

$$y = 1 \quad \text{if} \quad \sum_{i=1}^n w_i x_i + b \geq 0$$

# Mathematical Formulation: From Fixed to Flexible

## McCulloch-Pitts (The Old Way)

- **Inputs:** Binary  $x_i \in \{0, 1\}$
- **Weights:** None (all 1)
- **Threshold:** Fixed  $\theta$

$$y = 1 \quad \text{if} \quad \sum_{i=1}^n x_i \geq \theta$$

## The Perceptron (The New Way)

- **Inputs:** Real numbers  $x_i \in \mathbb{R}$
- **Weights:** Learnable  $w_i \in \mathbb{R}$
- **Bias:** Learnable  $b$  (or  $w_0$ )

$$y = 1 \quad \text{if} \quad \sum_{i=1}^n w_i x_i + b \geq 0$$


- We often rewrite the threshold  $\theta$  as a **bias term**  $b = -\theta$ .
- Moving it to the left side:  $\sum w_i x_i - \theta \geq 0 \implies \mathbf{w} \cdot \mathbf{x} + b \geq 0$ .
- **Compact Notation:**  $z = \mathbf{w}^T \mathbf{x} + b$ .

# Implementing Logic Gates with Weights

*By adjusting weights ( $w$ ) and bias ( $b$ ), the same architecture performs different tasks.*

## 1. The AND Gate

$y = 1$  if  $x_1 = 1, x_2 = 1$




$$1(1) + 1(1) - 1.5 = \mathbf{0.5} \geq 0 \implies 1$$

$$1(1) + 1(0) - 1.5 = -\mathbf{0.5} < 0 \implies 0$$

## 2. The OR Gate

$y = 1$  if any  $x = 1$



$$1(0) + 1(1) - 0.5 = \mathbf{0.5} \geq 0 \implies 1$$

$$1(0) + 1(0) - 0.5 = -\mathbf{0.5} < 0 \implies 0$$

# Real World Solution: The "Used Car" Revisited

**The Problem:** In the MP model, a Cup Holder ( $x_3$ ) was equal to Brakes ( $x_2$ ).

**The Perceptron Solution:** Assign **Weights** based on importance.

- $w_1$  (Engine) = 5
- $w_2$  (Brakes) = **10** (Critical!)
- $w_3$  (Cup Holder) = 1
- Bias  $b = -8$  (Threshold of 8)

# Real World Solution: The "Used Car" Revisited

**The Problem:** In the MP model, a Cup Holder ( $x_3$ ) was equal to Brakes ( $x_2$ ).

**The Perceptron Solution:** Assign **Weights** based on importance.

- $w_1$  (Engine) = 5
- $w_2$  (Brakes) = **10** (Critical!)
- $w_3$  (Cup Holder) = 1
- Bias  $b = -8$  (Threshold of 8)

## Scenario 1: No Brakes

Engine Good (1), **No Brakes (0)**, Cup Holder (1).

$$\sum = (5 \times 1) + (10 \times 0) + (1 \times 1) - 8$$

$$\sum = 6 - 8 = -2$$

$$-2 < 0 \implies \text{Don't Buy (0)}$$

# Real World Solution: The "Used Car" Revisited

**The Problem:** In the MP model, a Cup Holder ( $x_3$ ) was equal to Brakes ( $x_2$ ).

**The Perceptron Solution:** Assign **Weights** based on importance.

- $w_1$  (Engine) = 5
- $w_2$  (Brakes) = **10** (Critical!)
- $w_3$  (Cup Holder) = 1
- Bias  $b = -8$  (Threshold of 8)

## Scenario 1: No Brakes

Engine Good (1), **No Brakes (0)**, Cup Holder (1).

$$\sum = (5 \times 1) + (10 \times 0) + (1 \times 1) - 8$$

$$\sum = 6 - 8 = -2$$

$$-2 < 0 \implies \text{Don't Buy (0)}$$

## Scenario 2: Good Car

No Cup Holder (0), but Good Brakes (1).


$$\sum = (5 \times 1) + (10 \times 1) + (1 \times 0) - 8$$

$$\sum = 7$$

$$7 \geq 0 \implies \text{Buy (1)}$$




# Geometric Interpretation: The Hyperplane



## The Decision Boundary

- The equation  $\mathbf{w} \cdot \mathbf{x} + b = 0$  defines a **line** (in 2D) or a **hyperplane** (in  $n$ D).


# Geometric Interpretation: The Hyperplane



## The Decision Boundary

- The equation  $\mathbf{w} \cdot \mathbf{x} + b = 0$  defines a **line** (in 2D) or a **hyperplane** (in  $n$ D).
- This separates the input space into two regions:
  - $\mathbf{w} \cdot \mathbf{x} + b > 0 \implies$  Class 1
  - $\mathbf{w} \cdot \mathbf{x} + b < 0 \implies$  Class 0

# Geometric Interpretation: The Hyperplane



## The Decision Boundary

- The equation  $\mathbf{w} \cdot \mathbf{x} + b = 0$  defines a **line** (in 2D) or a **hyperplane** (in  $n$ D).
- This separates the input space into two regions:
  - $\mathbf{w} \cdot \mathbf{x} + b > 0 \implies$  Class 1
  - $\mathbf{w} \cdot \mathbf{x} + b < 0 \implies$  Class 0


## The Weight Vector ( $\mathbf{w}$ )

- The vector  $\mathbf{w}$  is always **orthogonal** (perpendicular) to the decision boundary.
- It points in the direction of the "Yes" (1) class.

# Geometric Intuition: Rotation vs. Translation

## 1. Changing Weights ( $w$ )

Controls the **Orientation**




### Rotation

Changing  $w$  **rotates** the boundary around the center. It changes *which feature matters more*.

## 2. Changing Bias ( $b$ )

Controls the **Position**



### Translation

Changing  $b$  **shifts** the boundary without rotating it. It changes the *threshold required to fire*.

# The Role of Bias ( $w_0$ ): Shifting the World

## Why do we need a Bias?

- Without a bias ( $b = 0$ ), the decision boundary **must pass through the origin**  $(0, 0)$ .
- The equation becomes  $\mathbf{w} \cdot \mathbf{x} = 0$ .
- **Problem:** What if the data isn't centered at the origin?


# The Role of Bias ( $w_0$ ): Shifting the World

## Why do we need a Bias?

- Without a bias ( $b = 0$ ), the decision boundary **must pass through the origin** ( $0, 0$ ).
- The equation becomes  $\mathbf{w} \cdot \mathbf{x} = 0$ .
- Problem:** What if the data isn't centered at the origin?

## The Solution:

- The Bias  $b$  acts as an "offset."
- It shifts the activation function left or right.
- It allows the decision boundary to move **off the origin** to fit the data better.



# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

Think of the **Bias** ( $b$  or  $w_0$ ) as your  
**Internal Disposition.**

- It represents how easy (or hard) it is to trigger you, *regardless of the evidence.*



# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

Think of the **Bias** ( $b$  or  $w_0$ ) as your  
**Internal Disposition.**

- It represents how easy (or hard) it is to trigger you, *regardless of the evidence.*

## The Golden Rule

$$\text{Evidence} + \text{Bias} > 0 \implies \text{Action}$$

# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

Think of the **Bias** ( $b$  or  $w_0$ ) as your  
**Internal Disposition.**

- It represents how easy (or hard) it is to trigger you, *regardless of the evidence.*

**Example: "Should I go to the Party?"**

## The Golden Rule

Evidence + Bias  $> 0 \implies$  Action

# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

Think of the **Bias** ( $b$  or  $w_0$ ) as your  
**Internal Disposition.**

- It represents how easy (or hard) it is to trigger you, *regardless of the evidence.*

## The Golden Rule

Evidence + Bias  $> 0 \implies$  Action

## Example: "Should I go to the Party?"

### The Extrovert (High Positive Bias)

**Bias** ( $b$ ) = +5

*"I love parties! I'm already halfway out the door."*

Even if the music is bad (Input  $x = -2$ ),  
 $-2 + 5 = 3 > 0 \implies$  **GO!**

# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

Think of the **Bias** ( $b$  or  $w_0$ ) as your  
**Internal Disposition.**

- It represents how easy (or hard) it is to trigger you, *regardless of the evidence.*

## The Golden Rule

Evidence + Bias  $> 0 \implies$  Action

## Example: "Should I go to the Party?"

### The Extrovert (High Positive Bias)

**Bias** ( $b$ ) = +5

*"I love parties! I'm already halfway out the door."*

Even if the music is bad (Input  $x = -2$ ),  
 $-2 + 5 = 3 > 0 \implies$  **GO!**

### The Homebody (High Negative Bias)

**Bias** ( $b$ ) = -5

*"I'd rather stay home. You need to convince me."*

The music must be AMAZING (Input  $x = +6$ ) just to get them to move,  $6 - 5 = 1 > 0 \implies$  **GO.**

# An Intuitive "Personality" Check

Think of the weighted sum ( $\sum w_i x_i$ ) as  
**External Evidence.**

- "Is the food good?"
- "Is it cheap?"

Think of the **Bias** ( $b$  or  $w_0$ ) as your  
**Internal Disposition.**

- It represents how easy (or hard) it is to trigger you, *regardless of the evidence.*

## The Golden Rule

Evidence + Bias  $> 0 \implies$  Action

*Bias adjusts the **threshold of activation**. It allows the neuron to say "Yes" easily or make it very strict.*

## Example: "Should I go to the Party?"

### The Extrovert (High Positive Bias)

**Bias** ( $b$ ) = +5

*"I love parties! I'm already halfway out the door."*

Even if the music is bad (Input  $x = -2$ ),  
 $-2 + 5 = 3 > 0 \implies$  **GO!**

### The Homebody (High Negative Bias)

**Bias** ( $b$ ) = -5

*"I'd rather stay home. You need to convince me."*

The music must be AMAZING (Input  $x = +6$ ) just to get them to move,  $6 - 5 = 1 > 0 \implies$  **GO.**

# The Learning Problem

# The Learning Problem

## The Goal:

Find a weight vector  $\mathbf{w}$  and bias  $b$  such that:

- For all Positive examples ( $x \in P$ ):  $\mathbf{w} \cdot \mathbf{x} + b \geq 0$
- For all Negative examples ( $x \in N$ ):  $\mathbf{w} \cdot \mathbf{x} + b < 0$

# The Learning Problem

## The Goal:

Find a weight vector  $\mathbf{w}$  and bias  $b$  such that:

- For all Positive examples ( $x \in P$ ):  $\mathbf{w} \cdot \mathbf{x} + b \geq 0$
- For all Negative examples ( $x \in N$ ):  $\mathbf{w} \cdot \mathbf{x} + b < 0$

## The Challenge:

We start with random weights. The line is wrong. How do we move it to the right place?



# The Learning Problem

## The Goal:

Find a weight vector  $\mathbf{w}$  and bias  $b$  such that:


- For all Positive examples ( $x \in P$ ):  $\mathbf{w} \cdot \mathbf{x} + b \geq 0$
- For all Negative examples ( $x \in N$ ):  $\mathbf{w} \cdot \mathbf{x} + b < 0$

## The Challenge:

We start with random weights. The line is wrong. How do we move it to the right place?

## The Approach:

*Iterative Error Correction.* We loop through the data, and every time the model makes a mistake, we "nudge" the weights.



# The Perceptron Learning Algorithm

**Input:** Training data  $D = \{(x, y)\}$

**Initialize:**  $\mathbf{w} \leftarrow \mathbf{0}$  (or random small numbers)

**Loop** until convergence (no errors/mistake):

For each pair  $(\mathbf{x}, y)$  in  $D$ :

① **Predict:**  $\hat{y} = \text{sign}(\mathbf{w} \cdot \mathbf{x})$

(1 if  $\geq 0$ , else 0)

② **Update:** If  $\hat{y} \neq y$  (Mistake!):

- If True is Positive ( $y = 1$ ) but predicted Negative:

$$\mathbf{w} \leftarrow \mathbf{w} + \mathbf{x}$$

- If True is Negative ( $y = 0$ ) but predicted Positive:

$$\mathbf{w} \leftarrow \mathbf{w} - \mathbf{x}$$

*Note: We absorb bias  $b$  into  $\mathbf{w}$  by adding a '1' to every input vector  $\mathbf{x}$ .*

# The Intuition: Steering the Weight Vector

## The Goal

We want the weight vector  $\mathbf{w}$  to point in the "correct" direction relative to our input data  $\mathbf{x}$ .

- For **Positive** examples ( $y = 1$ ),  $\mathbf{w}$  should point *roughly* in the same direction as  $\mathbf{x}$  (angle  $< 90^\circ$ ).
- For **Negative** examples ( $y = 0$ ),  $\mathbf{w}$  should point *away* from  $\mathbf{x}$  (angle  $> 90^\circ$ ).

**When the model makes a mistake, we simply "nudge"  $\mathbf{w}$  to fix the angle.**

**Too far away?**

Pull it closer (+)

**Too close?**

Push it away (−)


## Case 1: False Negative (The "Pull" Effect)

### The Error:

- Input  $\mathbf{x}$  is **Positive** ( $y = 1$ ).
- Model predicted **Negative** ( $\mathbf{w} \cdot \mathbf{x} < 0$ ).
- **Problem:** Angle is too wide ( $> 90^\circ$ ).

### The Fix: $\mathbf{w}_{new} \leftarrow \mathbf{w} + \mathbf{x}$

- We **add** the input vector to the weights.
- This pulls  **$\mathbf{w}$  towards  $\mathbf{x}$** .
- **Result:** Angle decreases, making  $\mathbf{w} \cdot \mathbf{x}$  more positive.



Angle Reduced!


## Case 2: False Positive (The "Push" Effect)

### The Error:

- Input  $\mathbf{x}$  is **Negative** ( $y = 0$ ).
- Model predicted **Positive** ( $\mathbf{w} \cdot \mathbf{x} > 0$ ).
- **Problem:** Angle is too narrow ( $< 90^\circ$ ).

### The Fix: $\mathbf{w}_{new} \leftarrow \mathbf{w} - \mathbf{x}$

- We **subtract** the input vector.
- This pushes  $\mathbf{w}$  **away** from  $\mathbf{x}$ .
- **Result:** Angle increases, making  $\mathbf{w} \cdot \mathbf{x}$  more negative.



Angle Increased!

# Mathematical Proof: Why the Dot Product Improves

Let  $\mathbf{w}_{new}$  be the updated weight vector. We check the new dot product  $\mathbf{w}_{new} \cdot \mathbf{x}$  to see if it moved in the right direction.

# Mathematical Proof: Why the Dot Product Improves

Let  $\mathbf{w}_{new}$  be the updated weight vector. We check the new dot product  $\mathbf{w}_{new} \cdot \mathbf{x}$  to see if it moved in the right direction.

## 1. False Negative Update ( $\mathbf{w} + \mathbf{x}$ )

We want the score to **increase** (become more positive).

$$\mathbf{w}_{new} \cdot \mathbf{x} = (\mathbf{w}_{old} + \mathbf{x}) \cdot \mathbf{x}$$

$$= \mathbf{w}_{old} \cdot \mathbf{x} + \mathbf{x} \cdot \mathbf{x}$$

$$= \underbrace{\mathbf{w}_{old} \cdot \mathbf{x}}_{\text{Old Score}} + \underbrace{\|\mathbf{x}\|^2}_{\text{Always Positive}}$$

**Conclusion:** The score **increases**.

# Mathematical Proof: Why the Dot Product Improves

Let  $\mathbf{w}_{new}$  be the updated weight vector. We check the new dot product  $\mathbf{w}_{new} \cdot \mathbf{x}$  to see if it moved in the right direction.

## 1. False Negative Update ( $\mathbf{w} + \mathbf{x}$ )

We want the score to **increase** (become more positive).

$$\begin{aligned}\mathbf{w}_{new} \cdot \mathbf{x} &= (\mathbf{w}_{old} + \mathbf{x}) \cdot \mathbf{x} \\ &= \mathbf{w}_{old} \cdot \mathbf{x} + \mathbf{x} \cdot \mathbf{x} \\ &= \underbrace{\mathbf{w}_{old} \cdot \mathbf{x}}_{\text{Old Score}} + \underbrace{\|\mathbf{x}\|^2}_{\text{Always Positive}}\end{aligned}$$

**Conclusion:** The score **increases**.

## 2. False Positive Update ( $\mathbf{w} - \mathbf{x}$ )

We want the score to **decrease** (become more negative).

$$\begin{aligned}\mathbf{w}_{new} \cdot \mathbf{x} &= (\mathbf{w}_{old} - \mathbf{x}) \cdot \mathbf{x} \\ &= \mathbf{w}_{old} \cdot \mathbf{x} - \mathbf{x} \cdot \mathbf{x} \\ &= \underbrace{\mathbf{w}_{old} \cdot \mathbf{x}}_{\text{Old Score}} - \underbrace{\|\mathbf{x}\|^2}_{\text{Always Positive}}\end{aligned}$$

**Conclusion:** The score **decreases**.



# Mathematical Proof: Why the Dot Product Improves

Let  $\mathbf{w}_{new}$  be the updated weight vector. We check the new dot product  $\mathbf{w}_{new} \cdot \mathbf{x}$  to see if it moved in the right direction.

## 1. False Negative Update ( $\mathbf{w} + \mathbf{x}$ )

We want the score to **increase** (become more positive).

$$\begin{aligned}\mathbf{w}_{new} \cdot \mathbf{x} &= (\mathbf{w}_{old} + \mathbf{x}) \cdot \mathbf{x} \\ &= \mathbf{w}_{old} \cdot \mathbf{x} + \mathbf{x} \cdot \mathbf{x} \\ &= \underbrace{\mathbf{w}_{old} \cdot \mathbf{x}}_{\text{Old Score}} + \underbrace{\|\mathbf{x}\|^2}_{\text{Always Positive}}\end{aligned}$$

**Conclusion:** The score **increases**.

## 2. False Positive Update ( $\mathbf{w} - \mathbf{x}$ )

We want the score to **decrease** (become more negative).

$$\begin{aligned}\mathbf{w}_{new} \cdot \mathbf{x} &= (\mathbf{w}_{old} - \mathbf{x}) \cdot \mathbf{x} \\ &= \mathbf{w}_{old} \cdot \mathbf{x} - \mathbf{x} \cdot \mathbf{x} \\ &= \underbrace{\mathbf{w}_{old} \cdot \mathbf{x}}_{\text{Old Score}} - \underbrace{\|\mathbf{x}\|^2}_{\text{Always Positive}}\end{aligned}$$

**Conclusion:** The score **decreases**.

*Since  $\mathbf{w} \cdot \mathbf{x} \propto \cos \theta$ , increasing the dot product decreases the angle  $\theta$ , and vice versa.*

# Real Inputs, Rigid Boundaries

## The Misconception

"Since we use real numbers  $(0.75, -1.2, \dots)$  instead of just 0 and 1, can't we solve complex problems?"

## The Misconception

"Since we use real numbers  $(0.75, -1.2, \dots)$  instead of just 0 and 1, can't we solve complex problems?"

**The Reality Check:** Recall the Perceptron equation:

$$\sum w_i x_i + b = 0$$

- This is the definition of a **Linear Hyperplane**.
- In 2D, it is a straight line ( $y = mx + c$ ).
- It **cannot bend**, curve, or encircle data.

# Real Inputs, Rigid Boundaries

## The Misconception



"Since we use real numbers  $(0.75, -1.2, \dots)$  instead of just 0 and 1, can't we solve complex problems?"

**The Reality Check:** Recall the Perceptron equation:

$$\sum w_i x_i + b = 0$$

- This is the definition of a **Linear Hyperplane**.
- In 2D, it is a straight line ( $y = mx + c$ ).
- It **cannot bend**, curve, or encircle data.

*Changing the domain from Binary to Real only fills the space; it doesn't bend the ruler.*



# Perceptron: Strengths & Limitations

## Strengths

# Perceptron: Strengths & Limitations

## Strengths

- **It Learns!**

An automatic learning rule adjusts weights from data.

## Strengths

- **It Learns!**

An automatic learning rule adjusts weights from data.

- **Feature Importance:**

Learns which inputs are more important via weights.

# Perceptron: Strengths & Limitations

## Strengths

- **It Learns!**  
An automatic learning rule adjusts weights from data.
- **Feature Importance:**  
Learns which inputs are more important via weights.
- **Convergence Guarantee:**  
If data is **linearly separable**, it's guaranteed to find a solution.



# Perceptron: Strengths & Limitations

## Strengths

- **It Learns!**  
An automatic learning rule adjusts weights from data.
- **Feature Importance:**  
Learns which inputs are more important via weights.
- **Convergence Guarantee:**  
If data is **linearly separable**, it's guaranteed to find a solution.

## Limitations

# Perceptron: Strengths & Limitations

## Strengths

- **It Learns!**  
An automatic learning rule adjusts weights from data.
- **Feature Importance:**  
Learns which inputs are more important via weights.
- **Convergence Guarantee:**  
If data is **linearly separable**, it's guaranteed to find a solution.

## Limitations

- **Linear Separability** Only problems where a single straight line (or hyperplane) can separate the classes can be solved.

# Perceptron: Strengths & Limitations

## Strengths

- **It Learns!**  
An automatic learning rule adjusts weights from data.
- **Feature Importance:**  
Learns which inputs are more important via weights.
- **Convergence Guarantee:**  
If data is **linearly separable**, it's guaranteed to find a solution.


## Limitations

- **Linear Separability** Only problems where a single straight line (or hyperplane) can separate the classes can be solved.
- **Harsh Threshold:** The step function activation is not differentiable, preventing modern gradient-based training.

# From Neuron to Network: Multilayer Perceptrons (MLP)

## 1. Input Layer

- The raw data ( $\mathbf{x}$ ).
- No computation happens here; it just passes values forward.




# From Neuron to Network: Multilayer Perceptrons (MLP)

## 1. Input Layer

- The raw data ( $\mathbf{x}$ ).
- No computation happens here; it just passes values forward.

## 2. Hidden Layer(s)

- The core innovation.
- These neurons are "hidden" from the outside world.
- They transform the input into **new features**.



# From Neuron to Network: Multilayer Perceptrons (MLP)

## 1. Input Layer


- The raw data ( $\mathbf{x}$ ).
- No computation happens here; it just passes values forward.

## 2. Hidden Layer(s)

- The core innovation.
- These neurons are "hidden" from the outside world.
- They transform the input into **new features**.

## 3. Output Layer

- The final decision ( $\hat{y}$ ).
- Combines the features from the



# How Hidden Layers Solve the Impossible

## The Logic: "Divide and Conquer"

A single Perceptron draws **one line**. It fails at XOR.


# How Hidden Layers Solve the Impossible

## The Logic: "Divide and Conquer"

A single Perceptron draws **one line**. It fails at XOR.

## A Network draws multiple lines:

- 1 **Hidden Neuron 1:** Draws Line A to separate the top-left.
- 2 **Hidden Neuron 2:** Draws Line B to separate the bottom-right.
- 3 **Output Neuron:** Combines them (e.g., "Active if Line A **OR** Line B says yes").



*Two lines creating a non-linear decision boundary.*




# How Hidden Layers Solve the Impossible

## The Logic: "Divide and Conquer"

A single Perceptron draws **one line**. It fails at XOR.

## A Network draws multiple lines:

- 1 **Hidden Neuron 1:** Draws Line A to separate the top-left.
- 2 **Hidden Neuron 2:** Draws Line B to separate the bottom-right.
- 3 **Output Neuron:** Combines them (e.g., "Active if Line A **OR** Line B says yes").



*Two lines creating a non-linear decision boundary.*

## Key Insight

Hidden layers transform the space so that the problem becomes linearly separable for the output layer.

# Visualizing the Transformation: XOR Solved

Let's use a simple MLP (one hidden layer and two neurons) with step activation functions (output is 0 or 1).

# Visualizing the Transformation: XOR Solved

Let's use a simple MLP (one hidden layer and two neurons) with step activation functions (output is 0 or 1).

- **Hidden 1 ( $h_1$ ):** Acts like OR ( $x_1 + x_2 \geq 0.5$ ) -  $w_0 = -0.5, w_1 = 1, w_2 = 1$

# Visualizing the Transformation: XOR Solved

Let's use a simple MLP (one hidden layer and two neurons) with step activation functions (output is 0 or 1).

- **Hidden 1 ( $h_1$ ):** Acts like OR ( $x_1 + x_2 \geq 0.5$ ) -  $w_0 = -0.5, w_1 = 1, w_2 = 1$
- **Hidden 2 ( $h_2$ ):** Acts like AND ( $x_1 + x_2 \geq 1.5$ ) -  $w_0 = -1.5, w_1 = 1, w_2 = 1$

# Visualizing the Transformation: XOR Solved

Let's use a simple MLP (one hidden layer and two neurons) with step activation functions (output is 0 or 1).

- **Hidden 1 ( $h_1$ ):** Acts like OR ( $x_1 + x_2 \geq 0.5$ ) -  $w_0 = -0.5, w_1 = 1, w_2 = 1$
- **Hidden 2 ( $h_2$ ):** Acts like AND ( $x_1 + x_2 \geq 1.5$ ) -  $w_0 = -1.5, w_1 = 1, w_2 = 1$
- **Output ( $y$ ):** Computes  $h_1 - h_2 \geq 0.5$  (essentially " $h_1$  AND NOT  $h_2$ ")

# Visualizing the Transformation: XOR Solved

Let's use a simple MLP (one hidden layer and two neurons) with step activation functions (output is 0 or 1).

- **Hidden 1 ( $h_1$ ):** Acts like OR ( $x_1 + x_2 \geq 0.5$ ) -  $w_0 = -0.5, w_1 = 1, w_2 = 1$
- **Hidden 2 ( $h_2$ ):** Acts like AND ( $x_1 + x_2 \geq 1.5$ ) -  $w_0 = -1.5, w_1 = 1, w_2 = 1$
- **Output ( $y$ ):** Computes  $h_1 - h_2 \geq 0.5$  (essentially " $h_1$  AND NOT  $h_2$ ")

## 1. The Computation Trace

Input Space		Hidden Space		Output
$x_1$	$x_2$	$h_1$	$h_2$	$y$
0	0	0	0	0
0	1	1	0	1
1	0	1	0	1
1	1	1	1	0

# Visualizing the Transformation: XOR Solved

Let's use a simple MLP (one hidden layer and two neurons) with step activation functions (output is 0 or 1).

- **Hidden 1 ( $h_1$ ):** Acts like OR ( $x_1 + x_2 \geq 0.5$ ) -  $w_0 = -0.5, w_1 = 1, w_2 = 1$
- **Hidden 2 ( $h_2$ ):** Acts like AND ( $x_1 + x_2 \geq 1.5$ ) -  $w_0 = -1.5, w_1 = 1, w_2 = 1$
- **Output ( $y$ ):** Computes  $h_1 - h_2 \geq 0.5$  (essentially " $h_1$  AND NOT  $h_2$ ")

## 1. The Computation Trace


Input Space		Hidden Space		Output
$x_1$	$x_2$	$h_1$	$h_2$	$y$
0	0	0	0	0
0	1	1	0	1
1	0	1	0	1
1	1	1	1	0

### Key Observation

Look at the inputs (0, 1) and (1, 0). They are distant in input space, but the hidden layer maps them to the **exact same point** (1, 0) in hidden space.

# Visualizing the Transformation: XOR Solved


## 2. The Geometric Transformation





# Visualizing the Transformation: XOR Solved


## 2. The Geometric Transformation




**A. Input Space (Inseparable)**

# Visualizing the Transformation: XOR Solved

## 2. The Geometric Transformation



A. Input Space (Inseparable)



B. Hidden Space (**Separable!**)

# Weights in a Network: Enter Linear Algebra

In a single perceptron, weights were a **vector** ( $\mathbf{w}$ ).

In a network, every layer is connected to the next.

- If Layer 1 has  $m$  neurons...
- And Layer 2 has  $n$  neurons...
- We need  $m \times n$  connections.

# Weights in a Network: Enter Linear Algebra

In a single perceptron, weights were a **vector** ( $\mathbf{w}$ ).

In a network, every layer is connected to the next.

- If Layer 1 has  $m$  neurons...
- And Layer 2 has  $n$  neurons...
- We need  $m \times n$  connections.

**The Weight Matrix  $\mathbf{W}$ :**

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

Where  $\mathbf{W}$  is a matrix of size  $(n \times m)$ .

- Row  $i$  contains the weights for Neuron  $i$  in the hidden layer.
- It learns to detect **one specific feature**.

# Weights in a Network: Enter Linear Algebra

In a single perceptron, weights were a **vector** ( $\mathbf{w}$ ).

In a network, every layer is connected to the next.


- If Layer 1 has  $m$  neurons...
- And Layer 2 has  $n$  neurons...
- We need  $m \times n$  connections.

**The Weight Matrix  $\mathbf{W}$ :**

$$\mathbf{h} = f(\mathbf{W}\mathbf{x} + \mathbf{b})$$

Where  $\mathbf{W}$  is a matrix of size  $(n \times m)$ .

- Row  $i$  contains the weights for Neuron  $i$  in the hidden layer.
- It learns to detect **one specific feature**.



# A Provocative Thought: "The lego block of intelligence?"

**Think about what we just did.**

- We used **2 neurons** to carve out a specific shape (XOR) by combining straight lines.
- We essentially created a "cut-out" in space.

# A Provocative Thought: "The lego block of intelligence?"

**Think about what we just did.**

- We used **2 neurons** to carve out a specific shape (XOR) by combining straight lines.
- We essentially created a "cut-out" in space.

## The Big Question

If 2 neurons can carve out an XOR shape... what could we carve out with **100 neurons**? Or **1,000,000**?


# A Provocative Thought: "The lego block of intelligence?"

**Think about what we just did.**

- We used **2 neurons** to carve out a specific shape (XOR) by combining straight lines.
- We essentially created a "cut-out" in space.

## The Big Question

If 2 neurons can carve out an XOR shape... what could we carve out with **100 neurons**? Or **1,000,000**?






# A Provocative Thought: "The lego block of intelligence?"

**Think about what we just did.**

- We used **2 neurons** to carve out a specific shape (XOR) by combining straight lines.
- We essentially created a "cut-out" in space.

## The Big Question

If 2 neurons can carve out an XOR shape... what could we carve out with **100 neurons**? Or **1,000,000**?



**The Hypothesis:** Maybe, just maybe, if we have enough neurons, we can approximate *any* shape, *any* decision boundary, or *any* mathematical function in the universe.

## Module 1B: Representation & Optimization

- **The Universal Approximation Theorem:**

Formalizing the idea that MLPs can approximate *any* continuous function.

- **The "Modern" Neuron:**

Why Step functions are dead ends for learning (Derivative is 0 or  $\infty$ ).

- **Enter Calculus:**

Transitioning to differentiable activation functions: **Sigmoid**, Tanh, and ReLU.

*Next Lecture: How do we teach the network automatically?*